

The Type Astronaut's Guide to Shapeless



November 2016

Dave Gurnell



The Type Astronaut's Guide to Shapeless

November 2016

Copyright 2016 Dave Gurnell. CC-BY-SA 3.0.

Published by Underscore Consulting LLP, Brighton, UK.

Contents

1	Introduction	7
1.1	What is generic programming?	7
1.2	About this book	9
1.3	Acknowledgements	11
I	Type Class Derivation	13
2	Algebraic data types and generic representations	15
2.1	Recap: algebraic data types	16
2.1.1	Alternative encodings	17
2.2	Generic product encodings	18
2.2.1	Switching representations using <i>Generic</i>	20
2.3	Generic coproducts	21
2.3.1	Switching encodings using <i>Generic</i>	22
2.4	Summary	23
3	Automatically deriving type class instances	25
3.1	Recap: type classes	25

3.1.1	Resolving instances	27
3.1.2	Idiomatic type class definitions	28
3.2	Deriving instances for products	30
3.2.1	Instances for <i>HLists</i>	31
3.2.2	Instances for concrete products	32
3.2.3	So what are the downsides?	34
3.3	Deriving instances for coproducts	36
3.3.1	Aligning CSV output	38
3.4	Deriving instances for recursive types	38
3.4.1	Implicit divergence	39
3.4.2	<i>Lazy</i>	40
3.5	Debugging implicit resolution	41
3.5.1	Debugging by summoning implicits	41
3.5.2	Debugging using <i>reify</i>	42
3.6	Summary	43
4	Working with types and implicits	45
4.1	Theory: dependent types	45
4.2	Dependently typed functions	47
4.3	Chaining dependent functions	51
4.4	Summary	54
5	Accessing names during implicit derivation	57
5.1	Literal types	57
5.2	Type tagging and phantom types	60
5.2.1	Records and <i>LabelledGeneric</i>	63

CONTENTS	5
5.3 Deriving product instances with <i>LabelledGeneric</i>	63
5.3.1 Instances for <i>HLists</i>	66
5.3.2 Instances for concrete products	68
5.4 Deriving coproduct instances with <i>LabelledGeneric</i>	69
5.5 Summary	71
II Shapeless Ops	73
6 Working with HLists and Coproducts	75
6.1 Simple ops examples	76
6.2 Creating a custom op	77
6.3 Case study: case class migrations	80
6.3.1 The type class	80
6.3.2 Step 1. Removing fields	81
6.3.3 Step 2. Reordering fields	83
6.3.4 Step 3. Adding new fields	84
6.4 Record ops	86
6.4.1 Selecting fields	87
6.4.2 Updating and removing fields	88
6.4.3 Converting to a regular Map	89
6.4.4 Other operations	89
6.5 Summary	89
7 Functional operations on HLists	91
7.1 Motivation: mapping over an HList	91
7.2 Polymorphic functions	92

7.2.1	How Polys work	92
7.2.2	Poly syntax	94
7.3	Mapping and flatMapping using Poly	97
7.4	Folding using Poly	99
7.5	Defining type classes using Poly	99
7.6	Summary	101
8	Counting with types	103
8.1	Representing numbers as types	103
8.2	Length of generic representations	104
8.3	Case study: random value generator	106
8.3.1	Simple random values	107
8.3.2	Random products	108
8.3.3	Random coproducts	109
8.4	Other operations involving Nat	111
8.5	Summary	112
	Preparing for launch	113
	TODOs	115

Chapter 1

Introduction

This book is a guide to using shapeless¹, a library for *generic programming* in Scala. Shapeless is large library, so rather than cover everything it has to offer we will concentrate on a few compelling use cases and use them to build a picture of the tools and patterns available.

Before we start, let's talk about what generic programming is and why shapeless is so exciting to Scala developers.

1.1 What is generic programming?

Types are helpful because they are specific: they show us how different pieces of code fit together, help us prevent bugs, and guide us toward solutions when we code.

Sometimes, however, types are *too* specific. There are situations where we want to exploit similarities between types to avoid repetition. For example, consider the following definitions:

¹<https://github.com/milessabin/shapeless>

```
case class Employee(name: String, number: Int, manager: Boolean)

case class IceCream(name: String, numCherries: Int, inCone: Boolean)
```

These two case classes represent different kinds of data but they have clear similarities: they both contain three fields of the same types. Suppose we want to implement a generic operation such as serializing to a CSV file. Despite the similarity between the two types, we have to write two separate serialization methods:

```
def employeeCsv(e: Employee): List[String] =
  List(e.name, e.number.toString, e.manager.toString)

def iceCreamCsv(c: IceCream): List[String] =
  List(c.name, c.numCherries.toString, c.inCone.toString)
```

Generic programming is about overcoming differences like these. Shapeless makes it convenient to convert specific types into generic ones that we can manipulate with common code.

For example, we can use the code below to convert employees and ice creams to values of the same type. Don't worry if you don't follow this example yet: we'll get to grips with the various concepts later on:

```
import shapeless._

val genericEmployee = Generic[Employee].to(Employee("Dave", 123, false))
// genericEmployee: shapeless::[String,shapeless::[Int,shapeless::[
  Boolean,shapeless.HNil]]] = Dave :: 123 :: false :: HNil

val genericIceCream = Generic[IceCream].to(IceCream("Sundae", 1, false))
// genericIceCream: shapeless::[String,shapeless::[Int,shapeless::[
  Boolean,shapeless.HNil]]] = Sundae :: 1 :: false :: HNil
```

Both values are now the same type. They are both heterogeneous lists (HLists for short) containing a `String`, an `Int`, and a `Boolean`. We will look at HLists and the important role they play soon. For now the important point is that both values are the same type. This means we can serialize them with the same function:


```
def genericCsv(gen: String :: Int :: Boolean :: HNil): List[String] =  
  List(gen(0), gen(1).toString, gen(2).toString)  
// genericCsv: (gen: shapeless.:::[String,shapeless.:::[Int,shapeless.:::[  
  Boolean,shapeless.HNil]]])List[String]  
  
genericCsv(genericEmployee)  
// res2: List[String] = List(Dave, 123, false)  
  
genericCsv(genericIceCream)  
// res3: List[String] = List(Sundae, 1, false)
```

This example is basic but it hints at the essence of generic programming. We reformulate problems so we can solve them use generic building blocks, and write small kernels of code that work with a wide variety of types. Generic programming with shapeless allows us to eliminate huge amounts of boilerplate, making Scala applications easier to read, write, and maintain.

Does that sound compelling? Thought so. Let's jump in!

1.2 About this book

This book is divided into two parts.

In Part I we introduce *type class derivation*, which allows us to create type class instances for any algebraic data type using only a handful of generic rules. Part I consists of four chapters:

- In Chapter 2 we introduce *generic representations*. We also introduce shapeless' Generic type class, which can produce a generic encoding for any case class or sealed trait.
- In Chapter 3 we use Generic to derive instances of a custom type class. We create an example type class to encode Scala data as Comma Separated Values (CSV), but the techniques we cover can be extended to many situations. We also introduce shapeless' Lazy type, which lets us handle recursive data like lists and trees.

- In Chapter 4 we introduce the theory and programming patterns we need to generalise the techniques from earlier chapters. Specifically we look at dependent types, dependently typed functions, and type level programming. This allows us to access more advanced applications of `shapeless`.
- In Chapter 5 we introduce `LabelledGeneric`, a variant of `Generic` that exposes field and type names as part of its generic representations. We also introduce additional theory: literal types, singleton types, phantom types, and type tagging. We demonstrate `LabelledGeneric` by creating a JSON encoder that preserves field and type names in its output.

In Part II we introduce the “ops type classes” provided in the `shapeless.ops` package. The ops type classes form a vast library of tools for manipulating generic representations. Rather than discuss every op in detail, we provide a theoretical primer in three chapters:

- In Chapter 6 we discuss the general layout of the ops type classes and provide an example that strings several simple ops together to form a powerful “case class migration” tool.
- In Chapter 7 we introduce *polymorphic functions* in `shapeless`, also known as `Polys`, and show how they are used in ops type classes for mapping, flat mapping, and folding over generic representations.
- Finally, in Chapter 8 we introduce the `Nat` type that `shapeless` uses to represent natural numbers at the type level. We introduce several related ops type classes, and use `Nat` to develop our own version of Scalacheck’s `Arbitrary`.

Source code and examples

This book is open source. See the inside cover for licensing information. You can find the Markdown source on Github:

<https://github.com/underscoreio/shapeless-guide>

The book receives constant updates from the community so be sure to

check the Github repo for the most up-to-date version.

We have also provided complete implementations of each of the major examples in an accompanying repo. See the README for details:

<https://github.com/underscoreio/shapeless-guide-code>

We assume shapeless 2.3.2 and either Typelevel Scala 2.11.8+ or Lightbend Scala 2.11.9+ / 2.12.1+.

1.3 Acknowledgements

Thanks to Miles Sabin, Richard Dallaway, Noel Welsh, Travis Brown, and our fellow space-farers on Github² for their invaluable contributions to this guide.

Special thanks to Sam Halliday for this excellent workshop Shapeless for Mortals³, which provided the initial inspiration and skeleton.

Finally, thanks to Rob Norris and his fellow contributors for the awesome Tut⁴, which keeps our examples compiling correctly.

²<https://github.com/underscoreio/shapeless-guide/graphs/contributors>

³<https://github.com/fommil/shapeless-for-mortals>

⁴<https://github.com/tpolecat/tut>

Part I

Type Class Derivation

Chapter 2

Algebraic data types and generic representations

The main idea behind generic programming is to solve problems for a wide variety of types by writing a small amount of generic code. Shapeless provides two sets of tools to this end:

1. a set of generic data types that can be inspected, traversed, and manipulated at the type level;
2. automatic mapping between *algebraic data types (ADTs)* (encoded in Scala as case classes and sealed traits) and these generic representations.

In this chapter we will start with a recap of the theory of algebraic data types and why they might be familiar to Scala developers. Then we will look at the data types shapeless uses as generic representations and discuss how they map on to concrete ADTs. Finally, we will introduce a type class called `Generic` that provides automatic mapping back and forth between ADTs and generic representations. We will finish with some simple examples using `Generic` to convert values from one type to another.

2.1 Recap: algebraic data types

Algebraic data types (ADTs¹) are a functional programming concept with a fancy name but a very simple meaning. They are an idiomatic way of representing data using “ands” and “ors”. For example:

- a shape is a rectangle **or** a circle
- a rectangle has a width **and** a height
- a circle has a radius

In ADT terminology, “and” types such as rectangle and circle are called *products* and “or” types such as shape are called *coproducts*. In Scala we typically represent products using case classes and coproducts using sealed traits:

```
sealed trait Shape
final case class Rectangle(width: Double, height: Double) extends Shape
final case class Circle(radius: Double) extends Shape

val rect: Shape = Rectangle(3.0, 4.0)
val circ: Shape = Circle(1.0)
```

The beauty of ADTs is that they are completely type safe. The compiler has complete knowledge of the algebras² we define, so it can help us write complete, correctly typed methods involving our types:

```
def area(shape: Shape): Double =
  shape match {
    case Rectangle(w, h) => w * h
    case Circle(r)       => math.Pi * r * r
  }

area(rect)
// res1: Double = 12.0
```

¹Be careful not to confuse algebraic data types with “abstract data types”, which are a different tool from computer science that has little bearing on the discussion here.

²The word “algebra” here means: the symbols we define, such as rectangle and circle; and the rules for manipulating those symbols, encoded as methods.


```
area(circ)
// res2: Double = 3.141592653589793
```

2.1.1 Alternative encodings

Sealed traits and case classes are undoubtedly the most convenient encoding of ADTs in Scala. However, they aren't the *only* encoding. For example, the Scala standard library provides generic products in the form of `Tuples` and a generic coproduct in the form of `Either`. We could have chosen these to encode our `Shape`:

```
type Rectangle2 = (Double, Double)
type Circle2    = Double
type Shape2     = Either[Rectangle2, Circle2]

val rect2: Shape2 = Left((3.0, 4.0))
val circ2: Shape2 = Right(1.0)
```

While this encoding is less readable than the case class encoding above, it does have some of the same desirable properties. We can still write completely type safe operations involving `Shape2`:

```
def area2(shape: Shape2): Double =
  shape match {
    case Left((w, h)) => w * h
    case Right(r)     => math.Pi * r * r
  }

area2(rect2)
// res4: Double = 12.0

area2(circ2)
// res5: Double = 3.141592653589793
```

Importantly, `Shape2` is a more *generic* encoding than `Shape`³. Any code that operates on a pair of `Doubles` will be able to operate on a `Rectangle2` and vice

³We're using "generic" with an informal way here rather than meaning "a type with a type parameter".

versa. As Scala developers we tend to prefer semantic types like `Rectangle` and `Circle` to generic ones like `Rectangle2` and `Circle2` precisely because of their specialised nature. However, in some cases generality is desirable. For example, if we're serializing data to disk, we don't care about the difference between a pair of `Doubles` and a `Rectangle2`. We just write two numbers and we're done.

`Shapeless` gives us the best of both worlds: we can use friendly semantic types by default and switch to generic representations when we want interoperability (more on this later). However, instead of using `Tuples` and `Either`, `shapeless` uses its own data types to represent generic products and coproducts. We'll introduce to these types in the next sections.

2.2 Generic product encodings

In the previous section we introduced tuples as a generic representation of products. Unfortunately, Scala's built-in tuples have a couple of disadvantages that make them unsuitable for `shapeless`' purposes:

1. Each size of tuple has a different, unrelated type, making it difficult to write code that abstracts over sizes.
2. There is no type for 0-length tuples, which are important for representing products with 0 fields. We could arguably use `Unit`, but we would ideally like all generic representations to have a common supertype. The least upper bound of `Unit` and `Tuple2` is `Any` so a combination of the two is not ideal.

For these reasons, `shapeless` uses a different generic encoding for product types called *heterogeneous lists* or `HLists`⁴.

An `HList` is either the empty list `HNil`, or a pair `:: [H, T]` where `H` is an arbitrary type and `T` is another `HList`. Because every `::` has its own `H` and `T`, the type of each element is encoded separately in the type of the overall list:

⁴Product is perhaps a better name for `HList`, but the standard library unfortunately already has a type `scala.Product`.

```
import shapeless.{HList, ::, HNil}

val product: String :: Int :: Boolean :: HNil =
  "Sunday" :: 1 :: false :: HNil
```

The type and value of the HList above mirror one another. Both represent three members: a String, an Int, and a Boolean. We can retrieve the head and tail and the types of the elements are preserved:

```
val first = product.head
// first: String = Sunday

val second = product.tail.head
// second: Int = 1

val rest = product.tail.tail
// rest: shapeless.::[Boolean,shapeless.HNil] = false :: HNil
```

The compiler knows the exact length of each HList, so it becomes a compilation error to take the head or tail of an empty list:

```
product.tail.tail.tail.head
// <console>:15: error: could not find implicit value for parameter c:
//     shapeless.ops.hlist.IsHCons[shapeless.HNil]
//     product.tail.tail.tail.head
//                               ^
```

We can manipulate and transform HLists in addition to being able to inspect and traverse them. For example, we can prepend an element with the `::` method. Again, notice how the type of the result reflects the number and types of its elements:

```
val newProduct: Long :: String :: Int :: Boolean :: HNil =
  42L :: product
```

Shapeless also provides tools for performing more complex operations such as mapping, filtering, and concatenating lists. We'll discuss these in more detail in Part II.

The behaviour described above isn't magic. We could have achieved all of this functionality using (A, B) and Unit as alternatives to :: and HNil. However, there is an advantage in keeping our representation types separate from the types used in our applications. HList provides this separation.

2.2.1 Switching representations using *Generic*

Shapeless provides a type class called *Generic* that allows us to switch back and forth between a concrete ADT and its generic representation. There's some macro magic going on behind the scenes that allows us to summon instances of *Generic* without boilerplate:

```
import shapeless.Generic

case class IceCream(name: String, numCherries: Int, inCone: Boolean)

val iceCreamGen = Generic[IceCream]
// iceCreamGen: shapeless.Generic[IceCream]{type Repr = shapeless.::[
  String,shapeless.::[Int,shapeless.::[Boolean,shapeless.HNil]]]} =
  anon$macro$4$l@2744a6f9
```

Note that the instance of *Generic* has a type member *Repr* containing the type of its generic representation. In this case *iceCreamGen.Repr* is *String :: Int :: Boolean :: HNil*. Instances of *Generic* have two methods: one for converting to the *Repr* type and one for converting from it:

```
val iceCream: IceCream =
  IceCream("Sundae", 1, false)
// iceCream: IceCream = IceCream(Sundae,1,false)

val repr: iceCreamGen.Repr =
  iceCreamGen.to(iceCream)
// repr: iceCreamGen.Repr = Sundae :: 1 :: false :: HNil

val iceCream2: IceCream =
  iceCreamGen.from(repr)
// iceCream2: IceCream = IceCream(Sundae,1,false)
```

If two ADTs have the same Repr, we can convert back and forth between them using their Generics:

```
case class Employee(name: String, number: Int, manager: Boolean)

// Create an employee from an ice cream:
val strangeEmployee: Employee =
  Generic[Employee].from(Generic[IceCream].to(iceCream))
// strangeEmployee: Employee = Employee(Sundae,1,false)
```

Other product types

It's useful to know that Generic understands tuples as well as case classes:

```
val tupleGen = Generic[(String, Int, Boolean)]

tupleGen.to(("Hello", 123, true))
// res4: tupleGen.Repr = Hello :: 123 :: true :: HNil

tupleGen.from("Hello" :: 123 :: true :: HNil)
// res5: (String, Int, Boolean) = (Hello,123,true)
```

2.3 Generic coproducts

Now we know how shapeless encodes product types. What about coproducts? We looked at Either earlier but that suffers from a similar drawback to tuples: we have no way of representing a disjunction of fewer than two types. For this reason, shapeless provides a different encoding that is similar to HList:

```
import shapeless.{Coproduct, :+:, CNil, Inl, Inr}

case class Red()
case class Amber()
case class Green()
```

```
type Light = Red :+: Amber :+: Green :+: CNil
```

In general coproducts take the form `A :+: B :+: C :+: CNil` meaning “A or B or C”, where `:+:` can be loosely interpreted as `Either`. The overall type of a coproduct encodes all the possible types in the disjunction, but each concrete instance contains a value for just one of the possibilities. `:+:` has two subtypes, `Inl` and `Inr`, that correspond loosely to `Left` and `Right`. We create instances of a coproduct by nesting `Inl` and `Inr` constructors:

```
val red    : Light = Inl(Red())
// red: Light = Inl(Red())

val green : Light = Inr(Inr(Inl(Green())))
// green: Light = Inr(Inr(Inl(Green())))
```

Every coproduct type is terminated with `CNil`, which is an empty type with no values, similar to `Nothing`.

We can never instantiate `CNil` or build a `Coproduct` purely from instances of `Inr`. We always have exactly one `Inl` in a value.

Again, it’s worth stating that `Coproducts` aren’t particularly special. The functionality above can be achieved using `Either` and `Nothing` in place of `:+:` and `CNil`. There are technical difficulties with using `Nothing`, but we could have used any other uninhabited or arbitrary singleton type in place of `CNil`.

2.3.1 Switching encodings using *Generic*

Coproduct types are difficult to parse on first glance. However, we can see how they fit into the larger picture of generic encodings. In addition to understanding case classes and case objects, shapeless’ `Generic` type class also understands sealed traits and abstract classes:

```
import shapeless.Generic

sealed trait Shape
final case class Rectangle(width: Double, height: Double) extends Shape
```

```
final case class Circle(radius: Double) extends Shape

val gen = Generic[Shape]
// gen: shapeless.Generic[Shape]{type Repr = shapeless.:+:[Rectangle,
    shapeless.:+:[Circle,shapeless.CNil]]} = anon$macro$1$1@69692800
```

The Repr of the Generic for Shape is a Coproduct of the subtypes of the sealed trait: `Rectangle :+: Circle :+: CNil`. We can use the `to` and `from` methods of the generic to map back and forth between Shape and `gen.Repr`:

```
gen.to(Rectangle(3.0, 4.0))
// res3: gen.Repr = Inl(Rectangle(3.0,4.0))

gen.to(Circle(1.0))
// res4: gen.Repr = Inr(Inl(Circle(1.0)))
```

2.4 Summary

In this chapter we discussed the generic representations shapeless provides for algebraic data types in Scala: `HLists` for product types and `Coproducts` for coproduct types. We also introduced the `Generic` type class to map back and forth between concrete ADTs and their generic representations. We haven't yet discussed why generic encodings are so attractive. The one use case we did cover—converting between ADTs—is fun but not tremendously useful.

The real power of `HLists` and `Coproducts` comes from their recursive structure. We can write code to traverse the structures and calculate values from their constituents. In the next chapter we will look at our first real use case: automatically deriving type class instances.

Chapter 3

Automatically deriving type class instances

In the last chapter we saw how the `Generic` type class allowed us to convert any instance of an ADT to a generic encoding made of `HLists` and `Coproducts`. In this chapter we will look at our first serious use case: automatic derivation of type class instances.

3.1 Recap: type classes

Before we get into the depths of instance derivation, let's quickly recap on the important aspects of type classes.

Type classes are a programming pattern borrowed from Haskell (the word “class” has nothing to do with classes in object oriented programming). We encode them in Scala using traits and implicits. A *type class* is a parameterised trait representing some sort of general functionality that we would like to apply to a wide range of types:

```
// Turn a value of type A into a row of cells in a CSV file:  
trait CsvEncoder[A] {
```

```
def encode(value: A): List[String]
}
```

We implement our type class with *instances* for each type we care about. If we want the instances to automatically be in scope we can place them in the type class' companion object. Otherwise we can place them in a separate library object for the user to import manually:

```
// Custom data type:
case class Employee(name: String, number: Int, manager: Boolean)

// CsvEncoder instance for the custom data type:
implicit val employeeEncoder: CsvEncoder[Employee] =
  new CsvEncoder[Employee] {
    def encode(e: Employee): List[String] =
      List(
        e.name,
        e.number.toString,
        if(e.manager) "yes" else "no"
      )
  }
```

We mark each instance with the keyword `implicit`, and define one or more *summoner* methods (also known as *materializers*) that accept an implicit parameter of the corresponding type:

```
def writeCsv[A](values: List[A])(implicit enc: CsvEncoder[A]): String =
  values.map(value => enc.encode(value).mkString(",")).mkString("\n")
```

When we call the entry point, the compiler calculates the value of the type parameter and searches for an implicit `CsvWriter` of the corresponding type:

```
val employees: List[Employee] = List(
  Employee("Bill", 1, true),
  Employee("Peter", 2, false),
  Employee("Milton", 3, false)
)

// The compiler inserts a CsvEncoder[Employee] parameter:
```

```
writeCsv(employees)
// res5: String =
// Bill,1,yes
// Peter,2,no
// Milton,3,no
```

We can use `writeCsv` with any data type we like, provided we have a corresponding implicit `CsvEncoder` in scope:

```
case class IceCream(name: String, numCherries: Int, inCone: Boolean)

implicit val iceCreamEncoder: CsvEncoder[IceCream] =
  new CsvEncoder[IceCream] {
    def encode(i: IceCream): List[String] =
      List(
        i.name,
        i.numCherries.toString,
        if(i.inCone) "yes" else "no"
      )
  }

val iceCreams: List[IceCream] = List(
  IceCream("Sundae", 1, false),
  IceCream("Cornetto", 0, true),
  IceCream("Banana Split", 0, false)
)

writeCsv(iceCreams)
// res8: String =
// Sundae,1,no
// Cornetto,0,yes
// Banana Split,0,no
```

3.1.1 Resolving instances

Type classes are very flexible but they require us to define instances for every type we care about. Fortunately, the Scala compiler has a few tricks up its sleeve to resolve instances for us given sets of user-defined rules. For example, we can write a rule that creates a `CsvEncoder` for `(A, B)` given `CsvEncoders` for `A` and `B`:

```
implicit def pairEncoder[A, B](
  implicit
  aEncoder: CsvEncoder[A],
  bEncoder: CsvEncoder[B]
): CsvEncoder[(A, B)] =
  new CsvEncoder[(A, B)] {
    def encode(pair: (A, B)): List[String] = {
      val (a, b) = pair
      aEncoder.encode(a) ++ bEncoder.encode(b)
    }
  }
}
```

When all the parameters to an `implicit def` are themselves marked as `implicit`, the compiler can use it as a resolution rule to create instances from other instances. For example, if we call `writeCsv` and pass in a `List[(Employee, IceCream)]`, the compiler is able to combine `pairEncoder`, `employeeEncoder`, and `iceCreamEncoder` to produce the required `CsvEncoder[(Employee, IceCream)]`:

```
writeCsv(employees zip iceCreams)
// res9: String =
// Bill,1,yes,Sundae,1,no
// Peter,2,no,Cornetto,0,yes
// Milton,3,no,Banana Split,0,no
```

Given a set of rules encoded as `implicit vals` and `implicit defs`, the compiler is capable of *searching* for combinations to give it the required instances. This behaviour, known as “implicit resolution”, is what makes the type class pattern so powerful in Scala.

Even with this power, the compiler can’t pull apart our case classes and sealed traits. We are required to define instances for ADTs by hand. Shapeless’ generic representations change all of this, allowing us to derive instances for any ADT for free.

3.1.2 Idiomatic type class definitions

The commonly accepted idiomatic style for type class definitions includes a companion object containing some standard methods:

```
object CsvEncoder {
  // "Summoner" method
  def apply[A](implicit enc: CsvEncoder[A]): CsvEncoder[A] =
    enc

  // "Constructor" method
  def instance[A](func: A => List[String]): CsvEncoder[A] =
    new CsvEncoder[A] {
      def encode(value: A): List[String] =
        func(value)
    }

  // Globally visible type class instances
}
```

The `apply` method, known as a “summoner” or “materializer”, allows us to summon a type class instance given a target type:

```
CsvEncoder[IceCream]
// res10: CsvEncoder[IceCream] = $anon$1@5916fe08
```

In simple cases the summoner does the same job as the `implicitly` method defined in `scala.Predef`:

```
implicitly[CsvEncoder[IceCream]]
// res11: CsvEncoder[IceCream] = $anon$1@5916fe08
```

As we will see in Section 4.2, when working with `shapeless` we encounter situations where `implicitly` doesn’t infer types correctly. We can always define the summoner method to do the right thing, so it’s worth writing one for every type class we create.

The `instance` method, sometimes named `pure`, provides a terse syntax for creating new type class instances, reducing the boilerplate of defining an anonymous class:

```
implicit val booleanEncoder: CsvEncoder[Boolean] =
  new CsvEncoder[Boolean] {
    def encode(b: Boolean): List[String] =
```

```
    if(b) List("yes") else List("no")
  }
```

down to something much shorter:

```
implicit val booleanEncoder: CsvEncoder[Boolean] =
  instance(b => if(b) List("yes") else List("no"))
```

Unfortunately, the layout of the code in this book prevents us writing long singletons containing lots of methods and instances. We therefore tend to define constructors and instances individually instead of in a companion object. Bear this in mind as you read and check the repo linked in Section 1.2 for complete worked examples.

3.2 Deriving instances for products

In this section we're going to use shapeless to derive type class instances for products (i.e. case classes). We'll use two intuitions:

1. If we have type class instances for the head and tail of an HList, we can derive an instance for the whole HList.
2. If we have a case class A, a `Generic[A]`, and a type class instance for the generic's Repr, we can combine them to create an instance for A.

Take `CsvEncoder` and `IceCream` as examples:

- `IceCream` has a generic Repr of type `String :: Int :: Boolean :: HNil`.
- The Repr is made up of a `String`, an `Int`, a `Boolean`, and an `HNil`. If we have `CsvEncoders` for these types, we can create an encoder for the whole thing.
- If we can derive a `CsvEncoder` for the Repr, we can create one for `IceCream`.

3.2.1 Instances for *HLists*

Let's start by writing `CsvEncoders`¹ for `String`, `Int`, and `Boolean`:

```
implicit val stringEncoder: CsvEncoder[String] =
  createEncoder(str => List(str))

implicit val intEncoder: CsvEncoder[Int] =
  createEncoder(num => List(num.toString))

implicit val booleanEncoder: CsvEncoder[Boolean] =
  createEncoder(bool => List(if(bool) "yes" else "no"))
```

We can combine these building blocks to create an encoder for our `HList`. We'll use two rules: one for `HNil` and one for `::` as shown below:

```
import shapeless.{HList, ::, HNil}

implicit val hnilEncoder: CsvEncoder[HNil] =
  createEncoder(hnil => Nil)

implicit def hlistEncoder[H, T <: HList](
  implicit
    hEncoder: CsvEncoder[H],
    tEncoder: CsvEncoder[T]
): CsvEncoder[H :: T] =
  createEncoder {
    case h :: t =>
      hEncoder.encode(h) ++ tEncoder.encode(t)
  }
```

Taken together, these five rules allow us to summon `CsvEncoders` for any `HList` involving `Strings`, `Ints`, and `Booleans`:

```
val reprEncoder: CsvEncoder[String :: Int :: Boolean :: HNil] =
  implicitly

reprEncoder.encode("abc" :: 123 :: true :: HNil)
```

¹Refer to Section 3.1 for the definition of `createEncoder`.

```
// res8: List[String] = List(abc, 123, yes)
```

3.2.2 Instances for concrete products

We can combine our derivation rules for HLists with an instance of `Generic` to produce a `CsvEncoder` for `IceCream`:

```
import shapeless.Generic

implicit val iceCreamEncoder: CsvEncoder[IceCream] = {
  val gen = Generic[IceCream]
  val enc = CsvEncoder[gen.Repr]
  createEncoder(iceCream => enc.encode(gen.to(iceCream)))
}
```

and use it as follows:

```
writeCsv(iceCreams)
// res10: String =
// Sundae,1,no
// Cornetto,0,yes
// Banana Split,0,no
```

This solution is specific to `IceCream`. Ideally we'd like to have a single rule that handles all case classes that have a `Generic` and a matching `CsvEncoder`. Let's work through the derivation step by step. Here's a first cut:

```
implicit def genericEncoder[A](
  implicit
  gen: Generic[A],
  enc: CsvEncoder[???]
): CsvEncoder[A] = createEncoder(a => enc.encode(gen.to(a)))
```

The first problem we have is selecting a type to put in place of the `???`. We want to write the `Repr` type associated with `gen`, but we can't do this:


```
implicit def genericEncoder[A](
  implicit
  gen: Generic[A],
  enc: CsvEncoder[gen.Repr]
): CsvEncoder[A] =
  createEncoder(a => enc.encode(gen.to(a)))
// <console>:27: error: illegal dependent method type: parameter may
//    only be referenced in a subsequent parameter section
//      gen: Generic[A],
//      ^
```

The problem here is a scoping issue: we can't refer to a type member of one parameter from another parameter in the same block. The trick to solving this is to introduce a new type parameter to our method and refer to it in each of the associated parameters:

```
implicit def genericEncoder[A, R](
  implicit
  gen: Generic[A] { type Repr = R },
  enc: CsvEncoder[R]
): CsvEncoder[A] =
  createEncoder(a => enc.encode(gen.to(a)))
```

We'll cover this coding style in more detail the next chapter. Suffice to say, this definition now compiles and works as expected and we can use it with any case class as expected. Intuitively, this definition says:

Given a type A and an HList type R, an implicit Generic to map A to R, and a CsvEncoder for R, create a CsvEncoder for A.

We now have a complete system that handles any case class. The compiler expands a call like:

```
writeCsv(iceCreams)
```

to use our family of derivation rules:

```
writeCsv(iceCreams)(  
  genericEncoder(  
    Generic[IceCream],  
    hlistEncoder(stringEncoder,  
      hlistEncoder(intEncoder,  
        hlistEncoder(booleanEncoder, hnilEncoder))))))
```

I'm sure you'll agree, it's nice not to have to write this code by hand!

Aux type aliases

Type refinements like `Generic[A] { type Repr = L }` are verbose and difficult to read, so shapeless provides a type alias `Generic.Aux` to rephrase the type member as a type parameter:

```
package shapeless  
  
object Generic {  
  type Aux[A, R] = Generic[A] { type Repr = R }  
}
```

Using this alias we get a much more readable definition:

```
implicit def genericEncoder[A, R](  
  implicit  
  gen: Generic.Aux[A, R],  
  env: CsvEncoder[R]  
): CsvEncoder[A] =  
  createEncoder(a => env.encode(gen.to(a)))
```

Note that the `Aux` type isn't changing any semantics—it's just making things easier to read. This pattern is used frequently in the shapeless codebase.

3.2.3 So what are the downsides?

If all of the above seems pretty magical, allow us to provide one significant dose of reality. If things go wrong, the compiler isn't great at telling us why.

There are two main reasons the code above might fail to compile. The first is when we can't find an implicit `Generic` instance. For example, here we try to call `writeCsv` with a non-case class:

```
class Foo(bar: String, baz: Int)

writeCsv(List(new Foo("abc", 123)))
// <console>:31: error: could not find implicit value for parameter
//      encoder: CsvEncoder[Foo]
//      writeCsv(List(new Foo("abc", 123)))
//              ^
```

In this case the error message is relatively easy to understand. If `shapeless` can't calculate a `Generic` it means that the type in question isn't an ADT—somewhere in the algebra there is a type that isn't a case class or a sealed abstract type.

The other potential source of failure is when the compiler can't calculate a `CsvEncoder` for our `HList`. This normally happens because we don't have an encoder for one of the fields in our ADT. For example, so far we haven't defined a `CsvEncoder` for `java.util.Date`, so the following code fails:

```
import java.util.Date

case class Booking(room: String, date: Date)

writeCsv(List(Booking("Lecture hall", new Date())))
// <console>:33: error: could not find implicit value for parameter
//      encoder: CsvEncoder[Booking]
//      writeCsv(List(Booking("Lecture hall", new Date())))
//              ^
```

The message we get here isn't very helpful. All the compiler knows is it tried a lot of combinations of implicits and couldn't make them work. It has no idea which combination came closest to the desired result, so it can't tell us where the source(s) of failure lie.

There's not much good news here. We have find the source of the error ourselves by a process of elimination. We'll discuss debugging techniques in Section 3.5. For now, the main redeeming feature is that implicit resolution always

fails at compile time. There's little chance that we will end up with code that fails during execution.

3.3 Deriving instances for coproducts

In the last section we created a set of rules to automatically derive a `CsvEncoder` for any product type. In this section we will apply the same patterns to coproducts. Let's return to our shape ADT as an example:

```
sealed trait Shape
final case class Rectangle(width: Double, height: Double) extends Shape
final case class Circle(radius: Double) extends Shape
```

The generic representation for Shape is `Rectangle :+: Circle :+: CNil`. In Section 3.2.2 we defined product encoders for `Rectangle` and `Circle`. Now, to write generic `CsvEncoders` for `:+:` and `CNil`, we can use the same principles we used for `HLists`:

```
import shapeless.{Coproduct, :+:, CNil, Inl, Inr}

implicit val cnilEncoder: CsvEncoder[CNil] =
  createEncoder(cnil => throw new Exception("Inconceivable!"))

implicit def coproductEncoder[H, T <: Coproduct](
  implicit
    hEncoder: CsvEncoder[H],
    tEncoder: CsvEncoder[T]
): CsvEncoder[H :+: T] = createEncoder {
  case Inl(h) => hEncoder.encode(h)
  case Inr(t) => tEncoder.encode(t)
}
```

There are two key points of note:

1. Alarming, the encoder for `CNil` throws an exception! Don't panic, though. Remember that we can't create values of type `CNil`. It's there as a marker for the compiler. It's ok to fail abruptly here because we will never reach this point. The `throw` expression is dead code.

2. Because Coproducts are *disjunctions* of types, the encoder for `:+:` has to *choose* whether to encode a left or right value. We pattern match on the two subtypes of `:+:`, which are `Inl` for left and `Inr` for right.

With these definitions and our product encoders from Section 3.2, we should be able to serialize a list of shapes. Let's give it a try:

```
val shapes: List[Shape] = List(
  Rectangle(3.0, 4.0),
  Circle(1.0)
)

writeCsv(shapes)
// <console>:33: error: could not find implicit value for parameter
//      encoder: CsvEncoder[Shape]
//      writeCsv(shapes)
//              ^
```

Oh no, it failed! The error message is unhelpful as we discussed earlier. The reason for the failure is we don't have a `CsvEncoder` instance for `Double`:

```
implicit val doubleEncoder: CsvEncoder[Double] =
  createEncoder(d => List(d.toString))
```

With this definition in place, everything works as expected:

```
writeCsv(shapes)
// res7: String =
// 3.0,4.0
// 1.0
```

SI-7046 and you

There is a Scala compiler bug called [SI-7046][link-si7046] that can cause coproduct generic resolution to fail. The bug causes certain parts of the macro API, on which shapeless depends, to be sensitive to the order in which types are defined in source code.

If you are using Lightbend Scala 2.11.8 or earlier and coproduct resolution fails for you, consider upgrading to Lightbend Scala 2.11.9 or Type-level Scala 2.11.8. SI-7046 is fixed in each of these releases.

3.3.1 Aligning CSV output

Our CSV encoder isn't very practical in its current form. It allows fields from `Rectangle` and `Circle` to occupy the same columns in the output. To fix this problem we need to modify the definition of `CsvEncoder` to incorporate the width of the data type and space the output accordingly. The examples repo linked in Section 1.2 contains a complete implementation of `CsvEncoder` that addresses this problem.

3.4 Deriving instances for recursive types

Let's try something more ambitious—a binary tree:

```
sealed trait Tree[A]
case class Branch[A](left: Tree[A], right: Tree[A]) extends Tree[A]
case class Leaf[A](value: A) extends Tree[A]
```

Theoretically we should already have all of the definitions in place to summon a CSV writer for this definition. However, calls to `writeCsv` fail to compile:

```
CsvEncoder[Tree[Int]]
// <console>:25: error: could not find implicit value for parameter enc:
//   CsvEncoder[Tree[Int]]
//       CsvEncoder[Tree[Int]]
//           ^
```

The problem is that our type is recursive. The compiler senses an infinite loop applying our derivation rules and gives up.

3.4.1 Implicit divergence

Implicit resolution is a search process. The compiler uses heuristics to determine whether it is “converging” on a solution. If the heuristics don’t yield favorable results for a particular branch of search, the compiler assumes the branch is not converging and moves onto another.

One heuristic is specifically designed to avoid infinite loops. If the compiler sees the same target type twice in a particular branch of search, it gives up and moves on. We can see this happening if we look at the expansion for `CsvEncode[Tree[Int]]`. The implicit resolution process goes through the following types:

```
CsvEncoder[Tree[Int]]           // 1
CsvEncoder[Branch[Int] :+: Leaf[Int] :+: CNil] // 2
CsvEncoder[Branch[Int]]         // 3
CsvEncoder[Tree[Int] :: Tree[Int] :: HNil]   // 4
CsvEncoder[Tree[Int]]           // 5 uh oh
```

We see `Tree[A]` twice in lines 1 and 5, so the compiler moves onto another branch of search. The result is failure to find a suitable implicit.

In fact, the situation is worse than this. If the compiler sees the same type constructor twice and the complexity of the type parameters is *increasing*, it assumes that branch of search is “diverging”. This is a problem for shapeless because types like `:: [H, T]` and `:+: [H, T]` come up in different generic representations and cause the compiler to give up prematurely. Consider the following types:

```
case class Bar(baz: Int, qux: String)
case class Foo(bar: Bar)
```

The expansion for `Foo` looks like this:

```
CsvEncoder[Foo]           // 1
CsvEncoder[Bar :: HNil]    // 2
```

```
CsvEncoder[Bar]           // 3
CsvEncoder[Int :: String :: HNil] // 4 uh oh
```

The compiler attempts to resolve a `CsvEncoder[::[H, T]]` twice in this branch of search, on lines 2 and 4. The type parameter for `T` is more complex on line 4 than on line 2, so the compiler assumes (incorrectly in this case) that the branch of search is diverging. It moves onto another branch and, again, the result is failure to find a suitable implicit.

3.4.2 *Lazy*

Implicit divergence would be a show-stopper for libraries like `shapeless`. Fortunately, `shapeless` provides a type called `Lazy` as a workaround. `Lazy` does two things:

1. it suppresses implicit divergence at compile time by guarding against some of the aforementioned over-defensive convergence heuristics;
2. it defers evaluation of the implicit parameter at runtime, permitting the derivation of self-referential implicits.

We use `Lazy` by wrapping it around specific implicit parameters. As a rule of thumb, it is always a good idea to wrap the “head” parameter of any `HList` or `Coproduct` rule and the `Repr` parameter of any `Generic` rule in `Lazy`:

```
implicit def hlistEncoder[H, T <: HList](
  implicit
    hEncoder: Lazy[CsvEncoder[H]], // wrapped in Lazy
    tEncoder: CsvEncoder[T]
): CsvEncoder[H :: T] = createEncoder {
  case h :: t =>
    hEncoder.value.encode(h) ++ tEncoder.encode(t)
}

implicit def coproductEncoder[H, T <: Coproduct](
  implicit
    hEncoder: Lazy[CsvEncoder[H]], // wrapped in Lazy
    tEncoder: CsvEncoder[T]
)
```



```

): CsvEncoder[H :+: T] = createEncoder {
  case Inl(h) => hEncoder.value.encode(h)
  case Inr(t) => tEncoder.encode(t)
}

implicit def genericEncoder[A, R](
  implicit
  gen: Generic.Aux[A, R],
  rEncoder: Lazy[CsvEncoder[R]] // wrapped in Lazy
): CsvEncoder[A] = createEncoder { value =>
  rEncoder.value.encode(gen.to(value))
}

```

This prevents the compiler giving up prematurely, and enables the solution to work on complex/recursive types like `Tree`:

```

CsvEncoder[Tree[Int]]
// res2: CsvEncoder[Tree[Int]] = $anon$1@300f9bbd

```

3.5 Debugging implicit resolution

Failures in implicit resolution can be confusing and frustrating. Here are a couple of techniques to use when implicits go bad.

3.5.1 Debugging by summoning implicits

What can we do when the compiler simply fails to find an implicit value? The failure could be caused by the resolution of any one of the implicits in use. For example:

```

case class Foo(bar: Int, baz: Float)

CsvEncoder[Foo]
// <console>:31: error: could not find implicit value for parameter enc:
   CsvEncoder[Foo]

```

```
//      CsvEncoder[Foo]
//      ^
```

The reason for the failure is that we haven't defined an `Ordering` for `Float`. However, this may not be obvious in application code. We can work through the expected expansion sequence to find the source of the error, inserting calls to `CsvEncoder.apply` or `implicitly` above the error to see if they compile. We start with the generic representation of `Foo`:

```
CsvEncoder[Int :: Float :: HNil]
// <console>:29: error: could not find implicit value for parameter enc:
//      CsvEncoder[shapeless.::[Int,shapeless.::[Float,shapeless.HNil]]]
//      ^
//      CsvEncoder[Int :: Float :: HNil]
```

This fails to compile so we know we can ignore `genericEncoder` for now. The next step is to try the components of the `HList`:

```
CsvEncoder[Int]

CsvEncoder[Float]
// <console>:29: error: could not find implicit value for parameter enc:
//      CsvEncoder[Float]
//      ^
//      CsvEncoder[Float]
```

`Int` passes but `Float` fails. `CsvEncoder[Float]` is a leaf in our tree of expansions, so we know to start by implementing this missing type class instance. If adding the instance doesn't fix the problem we repeat the process to find the next point of failure.

3.5.2 Debugging using `reify`

The `reify` method from `scala.reflect` takes a Scala expression as a parameter and returns an AST object representing the expression tree, complete with type annotations:

```
import scala.reflect.runtime.universe._

println(reify(CsvEncoder[Int]))
// Expr[CsvEncoder[Int]]($read.$iw.$iw.$iw.$iw.CsvEncoder.apply[Int](
//   $read.$iw.$iw.$iw.$iw.intEncoder))
```

The types inferred during implicit resolution can give us hints about problems. After implicit resolution, any remaining existential types such as `A` or `T` provide a sign that something has gone wrong. Similarly, “top” and “bottom” types such as `Any` and `Nothing` are evidence of failure.

3.6 Summary

In this chapter we discussed how to use `Generic`, `HLists`, and `Coproducts` to automatically derive type class instances. We also covered the `Lazy` type as a means of handling complex/recursive types. Taking all of this into account, we can write a common skeleton for deriving type class instances as follows.

First, define the type class:

```
import shapeless._

trait MyTC[A]
```

Define basic instances:

```
implicit def intInstance: MyTC[Int] = ???
implicit def stringInstance: MyTC[String] = ???
implicit def booleanInstance: MyTC[Boolean] = ???
```

Define instances for `HList`:

```
implicit def hnilInstance: MyTC[HNil] = ???

implicit def hlistInstance[H, T <: HList](
  implicit
    hInstance: Lazy[MyTC[H]], // wrap in Lazy
```

```
tInstance: MyTC[T]
): MyTC[H :: T] = ???
```

If required, define instances for Coproduct:

```
implicit def cnilInstance: MyTC[CNil] = ???

implicit def coproductInstance[H, T <: Coproduct](
  implicit
    hInstance: Lazy[MyTC[H]], // wrap in Lazy
    tInstance: MyTC[T]
): MyTC[H :+: T] = ???
```

Finally, define an instance for Generic:

```
implicit def genericInstance[A, R](
  implicit
    generic: Generic.Aux[A, R],
    rInstance: Lazy[MyTC[R]] // wrap in Lazy
): MyTC[A] = ???
```

In the next chapter we'll cover some useful theory, programming patterns, and debugging techniques to help write code in this style. In Chapter 5 we will revisit type class derivation using a variant of `Generic` that allows us to inspect field and type names in our ADTs.

Chapter 4

Working with types and implicits

In the last chapter we saw one of the most compelling use cases for shapeless: automatically deriving type class instances. There are plenty of even more powerful examples coming later. However, before we move on, we should take time to discuss some of the theory we've so far skipped over and establish a set of patterns for writing and debugging type- and implicit-heavy code.

4.1 Theory: dependent types

Last chapter we spent a lot of time using `Generic`, the type class for mapping ADT types to generic representations. However, we haven't yet discussed an important bit of theory that underpins `Generic` and much of shapeless: *dependent types*.

To illustrate this, let's take a closer look at `Generic`. Here's a simplified version of the definition:

```
package shapeless

trait Generic[A] {
  type Repr
  def to(value: A): Repr
```

```
def from(value: Repr): A
}
```

Instances of `Generic` reference two other types: a type parameter `A` and a type member `Repr`. Suppose we implement a method `getRepr` as follows. What type will we get back?

```
import shapeless.Generic

def getRepr[A](value: A)(implicit gen: Generic[A]) =
  gen.to(value)
```

The answer is it depends on the instance we get for `gen`. In expanding the call to `getRepr`, the compiler will search for a `Generic[A]` and the result type will be whatever `Repr` is defined in that instance:

```
case class Vec(x: Int, y: Int)
case class Rect(origin: Vec, size: Vec)

getRepr(Vec(1, 2))
// res1: shapeless.:[Int,shapeless.:[Int,shapeless.HNil]] = 1 :: 2 ::
//      HNil

getRepr(Rect(Vec(0, 0), Vec(5, 5)))
// res2: shapeless.:[Vec,shapeless.:[Vec,shapeless.HNil]] = Vec(0,0)
//      :: Vec(5,5) :: HNil
```

What we're seeing here is called *dependent typing*: the result type of `getRepr` depends on its value parameters via their type members. Suppose we had specified `Repr` as type parameter on `Generic` instead of a type member:

```
trait Generic2[A, Repr]

def getRepr2[A, R](value: A)(implicit generic: Generic2[A, R]): R =
  ???
```

We would have had to pass the desired value of `Repr` to `getRepr` as a type parameter, effectively making `getRepr` useless.

The intuitive take-away from this is that type parameters are useful as “inputs” and type members are useful as “outputs”.

4.2 Dependently typed functions

Shapeless uses dependent types all over the place: in `Generic`, in `Witness` (which we will see in the next chapter), and in a host of other implicit values that operate on `HLists`.

For example, shapeless provides a type class called `Last` that returns the last element in an `HList`. Here’s a simplified version of its definition:

```
package shapeless.ops.hlist

trait Last[L <: HList] {
  type Out
  def apply(in: L): Out =
    ??? // definition omitted for brevity
}
```

We can summon instances of `Last` to inspect `HLists` in our code. In the two examples below note that the `Out` types are dependent on the `HList` types we started with:

```
import shapeless.{HList, ::, HNil}

import shapeless.ops.hlist.Last

val last1 = Last[String :: Int :: HNil]
// last1: shapeless.ops.hlist.Last[shapeless.::[String,shapeless.::[Int,
  shapeless.HNil]]]{type Out = Int} = shapeless.ops.
  hlist$Last$$anon$34@4e65f7a0

val last2 = Last[Int :: String :: HNil]
// last2: shapeless.ops.hlist.Last[shapeless.::[Int,shapeless.::[String,
  shapeless.HNil]]]{type Out = String} = shapeless.ops.
  hlist$Last$$anon$34@7ca43724
```

Once we have summoned instances of `Last`, we can use them at the value level:

```
last1("foo" :: 123 :: HNil)
// res1: last1.Out = 123

last2(321 :: "bar" :: HNil)
// res2: last2.Out = bar
```

We get two forms of protection against errors. The implicits defined for `Last` ensure we can only summon instances if the input `HList` has at least one element:

```
Last[HNil]
// <console>:15: error: Implicit not found: shapeless.Ops.Last[shapeless
//      .HNil]. shapeless.HNil is empty, so there is no last element.
//      Last[HNil]
//      ^
```

In addition, the type parameters on the instances of `Last` check whether we pass in the expected type of `HList`:

```
last1(321 :: "bar" :: HNil)
// <console>:16: error: type mismatch;
// found   : shapeless.:[Int,shapeless.:[String,shapeless.HNil]]
// required: shapeless.:[String,shapeless.:[Int,shapeless.HNil]]
//      last1(321 :: "bar" :: HNil)
//      ^
```

As a further example, let's implement our own type class, called `Second`, that returns the second element in an `HList`:

```
trait Second[L <: HList] {
  type Out
  def apply(value: L): Out
}

object Second {
  type Aux[L <: HList, O] = Second[L] { type Out = O }
}
```



```
def apply[L <: HList](implicit inst: Second[L]): Aux[L, inst.Out] =
  inst
}
```

This code uses the idiomatic layout described in Section 3.1.2. We define the `Aux` type in the companion object beside the standard `apply` method for summoning instances.

Summoner methods versus “implicitly” versus “the”

Note that the return type on `apply` is `Aux[L, 0]`, not `Second[L]`. This is important. Using `Aux` prevents the `apply` method erasing the type members on summoned instances. If we define the return type as `Second[L]`, the `Out` type member will be erased from the return type and the type class will not work correctly.

The `implicitly` method from `scala.Predef` has this behaviour. Compare the type of an instance of `Last` summoned with `implicitly`:

```
implicitly[Last[String :: Int :: HNil]]
// res6: shapeless.ops.hlist.Last[shapeless...[String,shapeless
...:[Int,shapeless.HNil]]] = shapeless.ops.hlist$Last$$anon$34
@164e0044
```

to the type of an instance summoned with `Last.apply`:

```
Last[String :: Int :: HNil]
// res7: shapeless.ops.hlist.Last[shapeless...[String,shapeless
...:[Int,shapeless.HNil]]]{type Out = Int} = shapeless.ops.
hlist$Last$$anon$34@57b6c5f1
```

The type summoned by `implicitly` has no `Out` type member. For this reason, we should avoid `implicitly` when working with dependently typed functions. We can either use custom summoner methods, or we can use `shapeless`' replacement method, the:

```
import shapeless._

the[Last[String :: Int :: HNil]]
// res8: shapeless.ops.hlist.Last[shapeless...[String,shapeless
  ...:[Int,shapeless.HNil]]]{type Out = Int} = shapeless.ops.
  hlist$Last$$anon$34@c8acc98
```

We only need a single instance, defined for HLists of at least two elements:

```
implicit def hlistSecond[A, B, Rest <: HList]: Aux[A :: B :: Rest, B] =
  new Second[A :: B :: Rest] {
    type Out = B
    def apply(value: A :: B :: Rest): B =
      value.tail.head
  }
```

We can summon instances of `Second` subject to similar constraints to `Last`. If we try to summon an instance for an incompatible `HList`, resolution fails and we get a compile error:

```
val second1 = Second[String :: Boolean :: Int :: HNil]
// second1: Second[shapeless...[String,shapeless...[Boolean,shapeless
  ...:[Int,shapeless.HNil]]]]{type Out = Boolean} = $anon$1@504c9a5a

val second2 = Second[String :: Int :: Boolean :: HNil]
// second2: Second[shapeless...[String,shapeless...[Int,shapeless...[
  Boolean,shapeless.HNil]]]]{type Out = Int} = $anon$1@72fb9ba8

Second[String :: HNil]
// <console>:27: error: could not find implicit value for parameter inst
  : Second[shapeless...[String,shapeless.HNil]]
//       Second[String :: HNil]
//           ^
```

When we are able to summon an instance it comes with an `apply` method that works as expected at the value level:

```

second1("foo" :: true :: 123 :: HNil)
// res10: second1.Out = true

second2("bar" :: 321 :: false :: HNil)
// res11: second2.Out = 321

second1("baz" :: HNil)
// <console>:28: error: type mismatch;
// found   : shapeless.:[String,shapeless.HNil]
// required: shapeless.:[String,shapeless.:[Boolean,shapeless.:[Int,
//   shapeless.HNil]]]
//       second1("baz" :: HNil)
//                   ^

```

4.3 Chaining dependent functions

Dependently typed functions provide a means of calculating one type from another. We can *chain* dependently typed functions to perform calculations involving multiple steps. For example, we should be able to use a `Generic` to calculate a `Repr` for a case class, and use a `Last` to calculate the type of the last element. Let's try coding this:

```

def lastField[A](input: A)(
  implicit
    gen: Generic[A],
    last: Last[gen.Repr]
): last.Out = last.apply(gen.to(input))
// <console>:29: error: illegal dependent method type: parameter may
//   only be referenced in a subsequent parameter section
//       gen: Generic[A],
//           ^

```

Unfortunately our code doesn't compile. This is the same problem we had in Section 3.2.2 with our definition of `genericEncoder`. We worked around the problem by lifting the free type variable out as a type parameter:

```
def lastField[A, Repr <: HList](input: A)(
  implicit
    gen: Generic.Aux[A, Repr],
    last: Last[Repr]
): last.Out = last.apply(gen.to(input))

lastField(Rect(Vec(1, 2), Vec(3, 4)))
// res13: Vec = Vec(3,4)
```

As a general rule, we always write code in this style. By encoding all the free variables as type parameters, we enable the compiler to unify them with appropriate types. This goes for more subtle constraints as well. For example, suppose we wanted to summon a `Generic` for a case class of exactly one field. We might be tempted to write this:

```
def getWrappedValue[A, Head](input: A)(
  implicit
    gen: Generic.Aux[A, Head :: HNil]
): Head = gen.to(input).head
```

The result here is more insidious. The method definition compiles but the compiler can never find the implicits it needs at the call site:

```
case class Wrapper(value: Int)

getWrappedValue(Wrapper(42))
// <console>:31: error: could not find implicit value for parameter gen:
//      shapeless.Generic.Aux[Wrapper,shapeless.::[Head,shapeless.HNil]]
//      getWrappedValue(Wrapper(42))
//      ^
```

The error message hints at the problem:

```
error: could not find implicit value for parameter gen:
  Generic.Aux[Wrapper, Head :: HNil]
```

The clue is in the appearance of the type `Head`. This is the name of a type parameter in the method: it shouldn't be appearing in the type the compiler

is trying to unify. The problem is that the `gen` parameter is over-constrained: the compiler can't find a `Repr` and ensure its length at the same time. Nothing also often provides a clue, appearing when the compiler fails to unify covariant type parameters.

The solution to our problem above is to separate implicit resolution into steps:

1. find a `Generic` with a suitable `Repr` for `A`;
2. provide that the `Repr` has a `Head` type.

Here's a revised version of the method using `:=` to constrain `Repr`:

```
def getWrappedValue[A, Repr <: HList, Head, Tail <: HList](input: A)(
  implicit
    gen: Generic.Aux[A, Repr],
    ev: (Head :: Tail) := Repr
): Head = gen.to(input).head
// <console>:31: error: could not find implicit value for parameter c:
//      shapeless.ops.hlist.IsHCons[gen.Repr]
//      ): Head = gen.to(input).head
//      ^
```

This doesn't compile because the `head` method in the method body requires an implicit parameter of type `IsHCons`. This is a much simpler error message to fix—we just need to learn a tool from `shapeless`' toolbox. `IsHCons` is a `shapeless` type class that splits an `HList` into a `Head` and `Tail`. We can use `IsHCons` instead of `:=`:

```
import shapeless.ops.hlist.IsHCons

def getWrappedValue[A, Repr <: HList, Head, Tail <: HList](in: A)(
  implicit
    gen: Generic.Aux[A, Repr],
    isHCons: IsHCons.Aux[Repr, Head, Tail]
): Head = gen.to(in).head
```

This fixes the bug. Both the method definition and the call site now compile as expected:

```
getWrappedValue(wrapper(42))  
// res16: Int = 42
```

The take home point here isn't that we solved the problem using `IsHCons`. `Shapeless` provides a lot of tools like this (see Chapters 6 to 8), and we can supplement them where necessary with our own type classes. The important point is the process we used to write code that compiles and is capable of finding solutions. We'll finish off this section with a step-by-step guide summarising our findings so far.

4.4 Summary

When coding with `shapeless`, we are often trying to find a target type that depends on values in our code. This relationship is called *dependent typing*.

Problems involving dependent types can be conveniently expressed using implicit search, allowing the compiler to resolve intermediate and target types given a starting point at the call site.

We often have to use multiple steps to calculate a result (e.g. using a `Generic` to get a `Repr`, then using another type class to get to another type). When we do this, there are a few rules we can follow to ensure our code compiles and works as expected:

1. We should extract every intermediate type out to a type parameter. Many type parameters won't be used in the result, but the compiler needs them to know which types it has to unify.
2. The compiler resolves implicits from left to right, backtracking if it can't find a working combination. We should write implicits in the order we need them, using one or more type variables to connect them to previous implicits.
3. The compiler can only solve for one constraint at a time, so we mustn't over-constrain any single implicit.

4. We should state the return type explicitly, specifying any type parameters and type members that may be needed elsewhere. Type members are often important. If we don't state them in the return type, they won't be available to the compiler for further implicit resolution.
5. The Aux type alias pattern is useful for keeping code readable. We should look out for Aux aliases when using tools from the shapeless toolbox, and implement Aux aliases on our own dependently typed functions.

When we find a useful chain of dependently typed operations we can capture them as a single type class. This is sometimes called the “lemma” pattern (a term borrowed from mathematical proofs). We'll see an example of this pattern in Section 6.2.

Chapter 5

Accessing names during implicit derivation

Often, the type class instances we define need access to more than just types. In this chapter we will look at a variant of `Generic` called `LabelledGeneric` that gives us access to field names and type names.

To begin with we have some theory to cover. `LabelledGeneric` uses some clever techniques to expose name information at the type level. To understand these techniques we must discuss *literal types*, *singleton types*, *phantom types*, and *type tagging*.

5.1 Literal types

A Scala value may have multiple types. For example, the string "hello" has at least three types: `String`, `AnyRef`, and `Any`¹:

¹String has a bunch of other types like `Serializable` and `Comparable` but let's ignore those for now.

```
"hello" : String
// res0: String = hello

"hello" : AnyRef
// res1: AnyRef = hello

"hello" : Any
// res2: Any = hello
```

Interestingly, "hello" also has another type: a “singleton type” that belongs exclusively to that one value. This is similar to the singleton type we get when we define a companion object:

```
object Foo

Foo
// res3: Foo.type = Foo$@5d203f77
```

The type `Foo.type` is the type of `Foo`, and `Foo` is the only value with that type. Singleton types applied to literal values are called *literal types*. We don’t normally interact with them because the default behaviour of the compiler is to “widen” literals to their nearest non-singleton type. So, for example, these two expressions are essentially equivalent:

```
"hello"
// res4: String = hello

("hello" : String)
// res5: String = hello
```

Shapeless provides a few tools for working with literal types. First, there is a `narrow` macro that converts a literal expression into a singleton-typed literal expression:

```
import shapeless.syntax.singleton._
```

```
var x = 42.narrow
// x: Int(42) = 42
```

Note the type of `x` here: `Int(42)` is a literal type. It is a subtype of `Int` that only contains the value 42. If we attempt to assign a different number to `x`, we get a compile error:

```
x = 43
// <console>:16: error: type mismatch;
// found   : Int(43)
// required: Int(42)
//       x = 43
//           ^
```

However, `x` is still an `Int` according to normal subtyping rules. If we operate on `x` we get a regular type of result:

```
x + 1
// res6: Int = 43
```

We can use `narrow` on any literal in Scala:

```
1.narrow
// res7: Int(1) = 1

true.narrow
// res8: Boolean(true) = true

"hello".narrow
// res9: String("hello") = hello

// and so on...
```

However, we can't use it on compound expressions:

```
math.sqrt(4).narrow
// <console>:17: error: Expression scala.math.`package`.sqrt(4.0) does
//       not evaluate to a constant or a stable reference value
//       math.sqrt(4).narrow
```

```
//      ^  
// <console>:17: error: value narrow is not a member of Double  
//      math.sqrt(4).narrow  
//      ^
```

Literal types in Scala

Until recently, Scala had no syntax for writing literal types. The types were there in the compiler, but we couldn't express them directly in code. As of Lightbend Scala 2.12.1, Lightbend Scala 2.11.9, and Type-level Scala 2.11.8, however, we now have direct syntax support for literal types. In these versions of Scala we can write declarations like the following:

```
val theAnswer: 42 = 42
```

The type 42 is the same as the type `Int(42)` we saw in printed output earlier. You'll still see `Int(42)` in output for legacy reasons, but the canonical syntax going forward is 42.

5.2 Type tagging and phantom types

Shapeless uses literal types to model the names of fields in case classes. It does this by “tagging” the types of the fields with the literal types of their names. Before we see how shapeless does this, we'll do it ourselves to show that there's no magic (well... minimal magic, at any rate). Suppose we have a number:

```
val number = 42
```

This number is an `Int` in two worlds: at runtime, where it has an actual value and methods that we can call, and at compile-time, where the compiler uses the type to calculate which pieces of code work together and to search for implicits.

We can modify the type of number at compile time without modifying its run-time behaviour by “tagging” it with a “phantom type”. Phantom types are types with no run-time semantics, like this:

```
trait Cherries
```

We can tag number using `asInstanceOf`. We end up with a value that is both an `Int` and a `Cherries` at compile-time, and an `Int` at run-time:

```
val numCherries = number.asInstanceOf[Int with Cherries]
// numCherries: Int with Cherries = 42
```

Shapeless uses this trick to tag the types of fields in a case classes with the singleton types of their names. If you find using `asInstanceOf` uncomfortable then don't worry: shapeless provides two tagging syntaxes that avoid such unsavoriness.

The first syntax, `->>`, tags the expression on the right of the arrow with singleton type of the literal expression on the left:

```
import shapeless.labelled.{KeyTag, FieldType}
import shapeless.syntax.singleton._

val someNumber = 123

val numCherries = "numCherries" ->> someNumber
// numCherries: Int with shapeless.labelled.KeyTag[String("numCherries"), Int] = 123
```

Here we are tagging `someNumber` with the phantom type `KeyTag["numCherries", Int]`. The tag encodes both the name and type of the field, both of which are useful when searching for entries in a `Repr` using implicit resolution.

The second syntax takes the tag as a type rather than a literal value. This is useful when we know what tag to use but don't have the ability to write specific literals in our code:

```
import shapeless.labelled.field

field[Cherries](123)
```

Shapeless provides us with the `FieldType` type alias to make it easy to extract the key tag and value from a type:

```
type FieldType[K, V] = V with KeyTag[K, V]
```

As we'll see later, shapeless uses this mechanism to tag fields in products and subtypes in coproducts with identifiers from our source code. The tags exist purely at compile time and have no runtime representation.

If tags are just a phantom types how do we convert them to values we can use at runtime? Shapeless provides a type class called `Witness` for this purpose. If we combine `Witness` and `FieldType`, we get something very compelling—the ability extract the field name from a tagged field:

```
import shapeless.Witness

val numCherries = "numCherries" ->> 123
// numCherries: Int with shapeless.labelled.KeyTag[String("numCherries")
//              ,Int] = 123

// Get the tag from a tagged value:
def getFieldName[K, V](value: FieldType[K, V])
  (implicit witness: Witness.Aux[K]): K =
  witness.value

getFieldName(numCherries)
// res16: String = numCherries

// Get the untagged value from a tagged value:
def getFieldValue[K, V](value: FieldType[K, V]): V =
  value

getFieldValue(numCherries)
// res18: Int = 123
```

If we build an `HList` from values with tagged types we get a data structure that has some of the properties of a `Map`. We can reference fields by tag, manipulate

and replace them, and maintain all of the type and naming information along the way. Shapeless calls these structured “records”.

5.2.1 Records and *LabelledGeneric*

Records are HLists of items that are each tagged with type-level identifiers:

```
import shapeless.{HList, ::, HNil}

val garfield = ("cat" ->> "Garfield") :: ("orange" ->> true) :: HNil
// garfield: shapeless.::[String with shapeless.labelled.KeyTag[String("
    cat"),String],shapeless.::[Boolean with shapeless.labelled.KeyTag[
    String("orange"),Boolean],shapeless.HNil]] = Garfield :: true ::
    HNil
```

For clarity, the type of `garfield` is as follows:

```
// FieldType["cat", String] ::
// FieldType["orange", Boolean] ::
// HNil
```

We don't need to go into depth regarding records here; suffice to say that records are the generic representation used by the `LabelledGeneric` type class that we will discuss next. `LabelledGeneric` tags each item in a product or coproduct with the corresponding field or type name from the concrete ADT (although the names are represented as `Symbols`, not `Strings`). Shapeless provides a suite of Map-like operations on records, some of which we'll cover in Section 6.4. For now, though, let's derive some type classes using `LabelledGeneric`.

5.3 Deriving product instances with *LabelledGeneric*

We'll use a running example of JSON encoding to illustrate `LabelledGeneric`. We'll define a `JsonEncoder` type class that converts values to a JSON AST. This is the approach taken by Argonaut, Circe, Play JSON, Spray JSON, and many other Scala JSON libraries.

First we'll define our JSON data type:

```
sealed trait JsonValue
case class JsonObject(fields: List[(String, JsonValue)]) extends
  JsonValue
case class JsonArray(items: List[JsonValue]) extends JsonValue
case class JsonString(value: String) extends JsonValue
case class JsonNumber(value: Double) extends JsonValue
case class JsonBoolean(value: Boolean) extends JsonValue
case object JsonNull extends JsonValue
```

then the type class for encoding values as JSON:

```
trait JsonEncoder[A] {
  def encode(value: A): JsonValue
}

object JsonEncoder {
  def apply[A](implicit enc: JsonEncoder[A]): JsonEncoder[A] =
    enc
}
```

then a few basic instances:

```
def createEncoder[A](func: A => JsonValue): JsonEncoder[A] =
  new JsonEncoder[A] {
    def encode(value: A): JsonValue =
      func(value)
  }

implicit val stringEncoder: JsonEncoder[String] =
  createEncoder(str => JsonString(str))

implicit val doubleEncoder: JsonEncoder[Double] =
  createEncoder(num => JsonNumber(num))

implicit val intEncoder: JsonEncoder[Int] =
  createEncoder(num => JsonNumber(num))

implicit val booleanEncoder: JsonEncoder[Boolean] =
  createEncoder(bool => JsonBoolean(bool))
```



```
implicit def listEncoder[A]
  (implicit enc: JsonEncoder[A]): JsonEncoder[List[A]] =
  createEncoder(list => JsonArray(list.map(enc.encode)))

implicit def optionEncoder[A]
  (implicit enc: JsonEncoder[A]): JsonEncoder[Option[A]] =
  createEncoder(opt => opt.map(enc.encode).getOrElse(JsonNull))
```

Ideally, when we encode ADTs as JSON, we would like to use the correct field names in the output JSON:

```
case class IceCream(name: String, numCherries: Int, inCone: Boolean)

val iceCream = IceCream("Sundae", 1, false)

// Ideally we'd like to produce something like this:
val iceCreamJson: JsonValue =
  JsonObject(List(
    "name"      -> JsonString("Sundae"),
    "numCherries" -> JsonNumber(1),
    "inCone"     -> JsonBoolean(false)
  ))
```

This is where `LabelledGeneric` comes in. Let's summon an instance for `IceCream` and see what kind of representation it produces:

```
import shapeless.LabelledGeneric

val gen = LabelledGeneric[IceCream].to(iceCream)
// gen: shapeless.::[String with shapeless.labelled.KeyTag[Symbol with
  shapeless.tag.Tagged[String("name")],String],shapeless.::[Int with
  shapeless.labelled.KeyTag[Symbol with shapeless.tag.Tagged[String("
  numCherries")],Int],shapeless.::[Boolean with shapeless.labelled.
  KeyTag[Symbol with shapeless.tag.Tagged[String("inCone")],Boolean],
  shapeless.HNil]] = Sundae :: 1 :: false :: HNil
```

For clarity, the full type of the `HList` is:

```
// String with KeyTag[Symbol with Tagged["name"], String]    ::
// Int      with KeyTag[Symbol with Tagged["numCherries"], Int] ::
// Boolean with KeyTag[Symbol with Tagged["inCone"], Boolean]  ::
// HNil
```

The type here is slightly more complex than we have seen. Instead of representing the field names with literal string types, `shapeless` is representing them with symbols tagged with literal string types. The details of the implementation aren't particularly important: we can still use `Witness` and `FieldType` to extract the tags, but they come out as `Symbols` instead of `Strings`².

5.3.1 Instances for *HLists*

Let's define `JsonEncoder` instances for `HNil` and `::`. These encoders are going to generate and manipulate `JsonObject`s, so we'll introduce a new type of encoder to make that easier:

```
trait JsonObjectEncoder[A] extends JsonEncoder[A] {
  def encode(value: A): JsonObject
}

def createObjectEncoder[A](fn: A => JsonObject): JsonObjectEncoder[A] =
  new JsonObjectEncoder[A] {
    def encode(value: A): JsonObject =
      fn(value)
  }
```

The definition for `HNil` is then straightforward:

```
import shapeless.{HList, ::, HNil, Lazy}

implicit val hnilEncoder: JsonObjectEncoder[HNil] =
  createObjectEncoder(hnil => JsonObject(Nil))
```

The definition of `hlistEncoder` involves a few moving parts so we'll go through it piece by piece. We'll start with the definition we might expect if we were using regular `Generic`:

²Future versions of `shapeless` may switch to using `Strings` as tags instead of `Symbols`.

```
implicit def hlistObjectEncoder[H, T <: HList](
  implicit
    hEncoder: Lazy[JsonEncoder[H]],
    tEncoder: JsonObjectEncoder[T]
): JsonObjectEncoder[H :: T] = ???
```

LabelledGeneric will give us an HList of tagged types, so let's start by introducing a new type variable for the key type:

```
import shapeless.Witness
import shapeless.labelled.FieldType

implicit def hlistObjectEncoder[K, H, T <: HList](
  implicit
    hEncoder: Lazy[JsonEncoder[H]],
    tEncoder: JsonObjectEncoder[T]
): JsonObjectEncoder[FieldType[K, H] :: T] = ???
```

In the body of our method we're going to need the value associated with K. We'll add an implicit Witness[K] to do this for us:

```
implicit def hlistObjectEncoder[K, H, T <: HList](
  implicit
    witness: Witness.Aux[K],
    hEncoder: Lazy[JsonEncoder[H]],
    tEncoder: JsonObjectEncoder[T]
): JsonObjectEncoder[FieldType[K, H] :: T] = {
  val fieldName = witness.value
  ???
}
```

We can access the value of K using `witness.value`, but the compiler has no way of knowing what type of tag we're going to get. LabelledGeneric uses Symbols as the tag types, so we'll put a type bound on K and use `symbol.name` to convert it to a String:

```
implicit def hlistObjectEncoder[K <: Symbol, H, T <: HList](
  implicit
    witness: Witness.Aux[K],
```

```

hEncoder: Lazy[JsonEncoder[H]],
tEncoder: JsonObjectEncoder[T]
): JsonObjectEncoder[FieldType[K, H] :: T] = {
  val fieldName: String = witness.value.name
  ???
}

```

The rest of the definition uses the principles we covered in Chapter 3:

```

implicit def hlistObjectEncoder[K <: Symbol, H, T <: HList](
  implicit
    witness: Witness.Aux[K],
    hEncoder: Lazy[JsonEncoder[H]],
    tEncoder: JsonObjectEncoder[T]
): JsonObjectEncoder[FieldType[K, H] :: T] = {
  val fieldName: String = witness.value.name
  createObjectEncoder { hlist =>
    val head = hEncoder.value.encode(hlist.head)
    val tail = tEncoder.encode(hlist.tail)
    JsonObject((fieldName, head) :: tail.fields)
  }
}

```

5.3.2 Instances for concrete products

Finally let's turn to our generic instance. This is identical to the definitions we've seen before, except that we're using `LabelledGeneric` instead of `Generic`:

```

import shapeless.LabelledGeneric

implicit def genericObjectEncoder[A, H <: HList](
  implicit
    generic: LabelledGeneric.Aux[A, H],
    hEncoder: Lazy[JsonObjectEncoder[H]]
): JsonEncoder[A] =
  createObjectEncoder { value =>
    hEncoder.value.encode(generic.to(value))
  }
}

```

And that's all we need! With these definitions in place we can serialize instances of any case class and retain the field names in the resulting JSON:

```
JsonEncoder[IceCream].encode(iceCream)
// res15: JsonValue = JsonObject(List((name,JsonString(Sundae)), (
    numCherries,JsonNumber(1.0)), (inCone,JsonBoolean(false))))
```

5.4 Deriving coproduct instances with *LabelledGeneric*

Applying *LabelledGeneric* with Coproducts involves a mixture of the concepts we've covered already. Let's start by examining a Coproduct type derived by *LabelledGeneric*. We'll re-visit our Shape ADT from Chapter 3:

```
import shapeless.LabelledGeneric

sealed trait Shape
final case class Rectangle(width: Double, height: Double) extends Shape
final case class Circle(radius: Double) extends Shape

LabelledGeneric[Shape].to(Circle(1.0))
// res5: shapeless.:+:[Rectangle with shapeless.labelled.KeyTag[Symbol
    with shapeless.tag.Tagged[String("Rectangle")],Rectangle],shapeless
    .:+:[Circle with shapeless.labelled.KeyTag[Symbol with shapeless.
        tag.Tagged[String("Circle")],Circle],shapeless.CNil]] = Inr(Inl(
        Circle(1.0)))
```

Here is that Coproduct type in a more readable format:

```
// Rectangle with KeyTag[Symbol with Tagged["Rectangle"], Rectangle] :+:
// Circle    with KeyTag[Symbol with Tagged["Circle"],    Circle]    :+:
// CNil
```

As you can see, the result is a Coproduct of the subtypes of Shape, each tagged with the type name. We can use this information to write *JsonEncoders* for *:+:* and *CNil*:

```
import shapeless.{Coproduct, :+:, CNil, Inl, Inr, Witness, Lazy}
import shapeless.labelled.FieldType

implicit val cnilObjectEncoder: JsonObjectEncoder[CNil] =
  createObjectEncoder(cnil => throw new Exception("Inconceivable!"))

implicit def coproductObjectEncoder[K <: Symbol, H, T <: Coproduct](
  implicit
    witness: Witness.Aux[K],
    hEncoder: Lazy[JsonEncoder[H]],
    tEncoder: JsonObjectEncoder[T]
): JsonObjectEncoder[FieldType[K, H] :+: T] = {
  val typeName = witness.value.name
  createObjectEncoder {
    case Inl(h) =>
      JsonObject(List(typeName -> hEncoder.value.encode(h)))

    case Inr(t) =>
      tEncoder.encode(t)
  }
}
```

`coproductEncoder` follows the same pattern as `hlistEncoder`. We have three type parameters: `K` for the type name, `H` for the value at the head of the `HList`, and `T` for the value at the tail. We use `FieldType` and `:+:` in the result type to declare the relationships between the three, and we use a `Witness` to access the runtime value of the type name. The result is an object containing a single key/value pair, the key being the type name and the value the result:

```
val shape: Shape = Circle(1.0)

JsonEncoder[Shape].encode(shape)
// res8: JsonValue = JsonObject(List((Circle,JsonObject(List((radius,
  JsonNumber(1.0)))))))
```

5.5 Summary

In this chapter we discussed `LabelledGeneric`, a variant of `Generic` that exposes type and field names in its generic representations.

The names exposed by `LabelledGeneric` are encoded as type-level tags, so we can target them during implicit resolution. We started the chapter discussing *literal types* and the way `shapeless` uses them in its tags. We also discussed the `Witness` type class, which is used to reify literal types as values.

Finally, we brought `LabelledGeneric`, literal types, and `Witness` together to build a `JsonEncoder` library that includes sensible names in its output.

The key take home point from this chapter is that none of this code uses runtime reflection. Everything is implemented with types, implicits, and a small set of macros that are internal to `shapeless`. The code we're generating is consequently very fast and reliable at runtime.

Part II

Shapeless Ops

Chapter 6

Working with HLists and Coproducts

In Part I we discussed mechanisms for deriving type class instances for algebraic data types using `shapeless`. We can use these techniques to augment almost any type class, although in more complex cases we may have to write a lot of supporting code for manipulating HLists and Coproducts

In Part II we will look at the `shapeless.ops` package that provides a set of helpful tools that we can use as building blocks. Each op comes in two parts: a *type class* that we can use during implicit resolution, and *extension methods* that we can call on HList and Coproduct.

There are three general sets of ops, available from three packages:

- `shapeless.ops.hlist` defines type classes for HLists. Many can be used directly via extension methods on HList, defined in `shapeless.syntax.hlist`.
- `shapeless.ops.coproduct` defines type classes for Coproducts. Many can be used directly via extension methods on Coproduct, defined in `shapeless.syntax.coproduct`.
- `shapeless.ops.record` defines type classes for HLists whose elements are tagged with key types (Section 5.2). Many can be

used directly via extension methods on `HList`, defined in `shapeless.syntax.record`.

We don't have room in this book to cover all of the available ops. Fortunately, in most cases the code is understandable and well documented. Rather than provide an exhaustive guide, we will touch on the major theoretical and structural points and show you how to extract further information from the shapeless codebase.

6.1 Simple ops examples

`HList` has `init` and `last` extension methods based on two type classes: `shapeless.ops.hlist.Init` and `shapeless.ops.hlist.Last`. `Coproduct` has similar methods and type classes. These serve as perfect examples of the ops pattern. Here are simplified definitions of the extension methods:

```
package shapeless
package syntax

implicit class HListOps[L <: HList](l : L) {
  def last(implicit last: Last[L]): last.Out = last.apply(l)
  def init(implicit init: Init[L]): init.Out = init.apply(l)
}
```

The return type of each method is determined by a dependent type on the implicit parameter. The instances for each type class provide the actual mapping. Here's the skeleton definition of `Last` as an example:

```
trait Last[L <: HList] {
  type Out
  def apply(in: L): Out
}

object Last {
  type Aux[L <: HList, O] = Last[L] { type Out = O }
  implicit def pair[H]: Aux[H :: HNil, H] = ???
  implicit def list[H, T <: HList]
```

```
(implicit last: Last[T]): Aux[H :: T, last.Out] = ???
}
```

We can make a couple of interesting observations about this implementation. First, we can typically implement ops type classes with a small number of instances (just two in this case). We can therefore package *all* of the required instances in the companion object of the type class, allowing us to call the corresponding extension methods without any imports from `shapeless.ops`:

```
import shapeless._

("Hello" :: 123 :: true :: HNil).last
// res0: Boolean = true

("Hello" :: 123 :: true :: HNil).init
// res1: shapeless.::[String,shapeless.::[Int,shapeless.HNil]] = Hello
//      :: 123 :: HNil
```

Second, the type class is only defined for `HLists` with at least one element. This gives us a degree of static checking. If we try to call `last` on an empty `HList`, we get a compile error:

```
HNil.last
// <console>:16: error: Implicit not found: shapeless.Ops.Last[shapeless
//      .HNil.type]. shapeless.HNil.type is empty, so there is no last
//      element.
//      HNil.last
//      ^
```

6.2 Creating a custom op

Let's work through the creation of our own op as an exercise. We'll combine the power of `Last` and `Init` to create a `Penultimate` type class that retrieves the second-to-last element in an `HList`. This is an example of the “lemma” pattern we introduced in Section 4.4. Here's the type class definition, complete with `Aux` type alias and `apply` method:

```
import shapeless._

trait Penultimate[L] {
  type Out
  def apply(l: L): Out
}

object Penultimate {
  type Aux[L, O] = Penultimate[L] { type Out = O }

  def apply[L](implicit p: Penultimate[L]): Aux[L, p.Out] = p
}
```

We only need to define one instance, combining `Init` and `Last` using the techniques covered in Section 4.3:

```
import shapeless.ops.hlist

implicit def hlistPenultimate[L <: HList, M <: HList, O](
  implicit
    init: hlist.Init.Aux[L, M],
    last: hlist.Last.Aux[M, O]
): Penultimate.Aux[L, O] =
  new Penultimate[L] {
    type Out = O
    def apply(l: L): O =
      last.apply(init.apply(l))
  }
```

We can use `Penultimate` as follows:

```
type BigList =
  String :: Int :: Boolean :: Double :: HNil

val bigList: BigList =
  "foo" :: 123 :: true :: 456.0 :: HNil

Penultimate[BigList].apply(bigList)
// res4: Boolean = true
```

Summoning an instance of `Penultimate` requires the compiler to summon

instances for Last and Init, so we inherit the same level of type checking on short HLists:

```
type TinyList = String :: HNil

val tinyList = "bar" :: HNil

Penultimate[TinyList].apply(tinyList)
// <console>:22: error: could not find implicit value for parameter p:
//     Penultimate[TinyList]
//           Penultimate[TinyList].apply(tinyList)
//                                     ^
```

We can make things more convenient for end users by defining an extension method on HList:

```
implicit class PenultimateOps[A](a: A) {
  def penultimate(implicit inst: Penultimate[A]): inst.Out =
    inst.apply(a)
}

bigList.penultimate
// res7: Boolean = true
```

We can also provide Penultimate for all product types by providing an instance based on Generic:

```
implicit def genericPenultimate[A, R, O](
  implicit
    generic: Generic.Aux[A, R],
    penultimate: Penultimate.Aux[R, O]
): Penultimate.Aux[A, O] =
  new Penultimate[A] {
    type Out = O
    def apply(a: A): O =
      penultimate.apply(generic.to(a))
  }

case class IceCream(name: String, numCherries: Int, inCone: Boolean)

IceCream("Sundae", 1, false).penultimate
```

```
// res9: Int = 1
```

6.3 Case study: case class migrations

The power of ops type classes fully crystallizes when we chain them together as building blocks for our own code. We'll finish this chapter with a powerful example: a type class for performing “migrations” (aka “evolutions”) on case classes¹. For example, if version 1 of our app has a case class:

```
case class IceCreamV1(name: String, numCherries: Int, inCone: Boolean)
```

our type class should let us perform certain mechanical “upgrades” for free:

```
// Remove fields:
case class IceCreamV2a(name: String, inCone: Boolean)

// Reorder fields:
case class IceCreamV2b(name: String, inCone: Boolean, numCherries: Int)

// Insert fields (provided we can determine a default value):
case class IceCreamV2c(
  name: String, inCone: Boolean, numCherries: Int, numWaffles: Int)
```

Ideally we'd like to be able to write code like this:

```
IceCreamV1("Sundae", 1, false).migrateTo[IceCreamV2a]
```

The type class should take care of the migration without additional boilerplate.

6.3.1 The type class

The Migration type class represents a transformation from a source to a destination type. Both of these are going to be “input” types in our derivation, so

¹The term is stolen from “database migrations”—SQL scripts that automate upgrades to a database schema.

we model both as type parameters. We don't need an Aux type alias because there are no type members to expose:

```
trait Migration[A, B] {
  def apply(a: A): B
}
```

We'll introduce an extension method now to make later examples easier to read:

```
implicit class MigrationOps[A](a: A) {
  def migrateTo[B](implicit migration: Migration[A, B]): B =
    migration.apply(a)
}
```

6.3.2 Step 1. Removing fields

Let's build up the solution piece by piece, starting with field removal. We can do this in several steps:

1. convert A to its generic representation;
2. filter the HList from step 1—only retain fields that are also in B;
3. convert the output of step 2 to B.

We can implement steps 1 and 3 with `Generic` or `LabelledGeneric`, and step 2 with an op called `Intersection`:

```
import shapeless._
import shapeless.ops.hlist

implicit def genericMigration[A, B, ARepr <: HList, BRepr <: HList](
  implicit
    aGen    : Generic.Aux[A, ARepr],
    bGen    : Generic.Aux[B, BRepr],
    inter   : hlist.Intersection.Aux[ARepr, BRepr, BRepr]
): Migration[A, B] = new Migration[A, B] {
  def apply(a: A): B =
    bGen.from(inter.apply(aGen.to(a)))
}
```

```
}
```

Take a moment to locate `Intersection`² in the shapeless codebase. Its `Aux` type alias takes three parameters: two input `HLists` and one output for the intersection type. In the example above we are specifying `ARepr` and `BRepr` as the input types and `BRepr` as the output type. This means implicit resolution will only succeed if `B` has an exact subset of the fields of `A`, specified with the exact same names in the same order:

```
IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2a]
// res6: IceCreamV2a = IceCreamV2a(Sundae,true)
```

We get a compile error if we try to use `Migration` with non-conforming types:

```
IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2b]
// <console>:24: error: could not find implicit value for parameter
//      migration: Migration[IceCreamV1,IceCreamV2b]
//      IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2b]
//                                     ^
```

Field uniqueness

`Intersection` doesn't use the field names from `LabelledGeneric`. What would have happened if we had used `Generic` above instead?

Although we're not using the field names directly, they are having an effect on the uniqueness checks `Intersection` uses to identify fields. If we had used regular `Generic`, the behaviour would be different and likely incorrect.

²<https://github.com/milessabin/shapeless/blob/shapeless-2.3.2/core/src/main/scala/shapeless/ops/hlists.scala#L1297-L1352>

6.3.3 Step 2. Reordering fields

We need to lean on another ops type class to add support for reordering. `hlist.Align`³ allows us reorder the fields in one `HList` to match the order they appear in another `HList`. We can redefine our instance using `Align` as follows:

```
implicit def genericMigration[
  A, B,
  ARepr <: HList, BRepr <: HList,
  Unaligned <: HList
](
  implicit
    aGen  : LabelledGeneric.Aux[A, ARepr],
    bGen  : LabelledGeneric.Aux[B, BRepr],
    inter : hlist.Intersection.Aux[ARepr, BRepr, Unaligned],
    align : hlist.Align[Unaligned, BRepr]
): Migration[A, B] = new Migration[A, B] {
  def apply(a: A): B =
    bGen.from(align.apply(inter.apply(aGen.to(a))))
}
```

We introduce a new type parameter called `Unaligned` to represent the intersection of `ARepr` and `BRepr` before alignment, and use `Align` to convert `Unaligned` to `BRepr`. With this modified type class instance we can support the removal and reordering of fields:

```
IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2a]
// res8: IceCreamV2a = IceCreamV2a(Sundae,true)

IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2b]
// res9: IceCreamV2b = IceCreamV2b(Sundae,true,1)
```

However, we still get a failure if we try to add fields:

³<https://github.com/milessabin/shapeless/blob/shapeless-2.3.2/core/src/main/scala/shapeless/ops/hlists.scala#L1973-L1997>

```
IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2c]
// <console>:26: error: could not find implicit value for parameter
//      migration: Migration[IceCreamV1,IceCreamV2c]
//      IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2c]
//                                         ^
```

6.3.4 Step 3. Adding new fields

We need a mechanism for calculating default values to suppose the addition of new fields. Shapeless doesn't provide a type class for this but Cats does in the form of a `Monoid`. Here's a simplified definition:

```
package cats

trait Monoid[A] {
  def empty: A
  def combine(x: A, y: A): A
}
```

`Monoid` defines two operations: `empty` for creating a “default” value of the type and `combine` for “adding” two values. We don't need `combine` in our code but in most cases it is trivial to define.

Cats provides instance of `Monoid` for all of the primitive types we care about (`Int`, `Double`, `Boolean`, and `String`). We can define instances for `HNil` and `::` using the techniques from Chapter 5:

```
import cats.Monoid
import cats.instances.all._
import shapeless.labelled.{field, FieldType}

def createMonoid[A](zero: A)(add: (A, A) => A): Monoid[A] =
  new Monoid[A] {
    def empty = zero
    def combine(x: A, y: A): A = add(x, y)
  }

implicit val hnilMonoid: Monoid[HNil] =
  createMonoid[HNil](HNil)((x, y) => HNil)
```

```
implicit def emptyHList[K <: Symbol, H, T <: HList](
  implicit
    hMonoid: Lazy[Monoid[H]],
    tMonoid: Monoid[T]
): Monoid[FieldType[K, H] :: T] =
  createMonoid(field[K](hMonoid.value.empty) :: tMonoid.empty) {
    (x, y) =>
      field[K](hMonoid.value.combine(x.head, y.head)) ::
        tMonoid.combine(x.tail, y.tail)
  }
```

We need to combine `Monoid`⁴ with a couple of other ops to complete our final implementation of `Migration`. Here's the full list of steps:

1. use `LabelledGeneric` to convert A to its generic representation;
2. use `Intersection` to calculate an `HList` of fields common to A and B;
3. calculate the types of fields that appear in B but not in A;
4. use `Monoid` to calculate a default value of the type from step 3;
5. append the common fields from step 2 to the new field from step 4;
6. use `Align` to reorder the fields from step 5 in the same order as B;
7. use `LabelledGeneric` to convert the output of step 6 to B.

We've already seen how to implement steps 1, 2, 4, 6, and 7. We can implement step 3 using an ops type class called `Diff` that works very similarly to `Intersection`. We can implement step 5 using another ops type class called `Prepend`. Here's the complete solution:

```
implicit def genericMigration[
  A, B, ARepr <: HList, BRepr <: HList,
  Common <: HList, Added <: HList, Unaligned <: HList
](
  implicit
    aGen : LabelledGeneric.Aux[A, ARepr],
    bGen : LabelledGeneric.Aux[B, BRepr],
    inter : hlist.Intersection.Aux[ARepr, BRepr, Common],
    diff : hlist.Diff.Aux[BRepr, Common, Added],
```

⁴Pun intended.

```

monoid : Monoid[Added],
prepend : hlist.Prepend.Aux[Added, Common, Unaligned],
align   : hlist.Align[Unaligned, BRepr]
): Migration[A, B] =
  new Migration[A, B] {
    def apply(a: A): B =
      bGen.from(align(prepend(monoid.empty, inter(aGen.to(a)))))
  }

```

Note that we don't end up using every type class at the value level. We use `Diff` to calculate the `Added` data type, but we don't actually need `diff.apply` at run time. Instead we use `Monoid` to summon an instance of `Added`.

With this final version of the type class instance in place we can use `Migration` for all the use cases we set out at the beginning of the case study:

```

IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2a]
// res14: IceCreamV2a = IceCreamV2a(Sundae,true)

IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2b]
// res15: IceCreamV2b = IceCreamV2b(Sundae,true,1)

IceCreamV1("Sundae", 1, true).migrateTo[IceCreamV2c]
// res16: IceCreamV2c = IceCreamV2c(Sundae,true,1,0)

```

It's amazing what we can create with ops type classes. `Migration` has a single implicit `def` with a single line of value-level implementation. It allows us to automate migrations between *any* pair of case classes, in roughly the same amount of code we'd write to handle a *single* pair of types using the standard library. It's a wonderful demonstration of the power of shapeless!

6.4 Record ops

We've spent some time in this chapter looking at type classes from the `shapeless.ops.hlist` and `shapeless.ops.coproduct` packages. We mustn't leave without mentioning a third important package: `shapeless.ops.record`.

Shapeless' "record ops" provide Map-like operations on HLists of tagged elements. Here are a handful of examples of the many available ops. Our examples will involve ice creams:

```
import shapeless._

case class IceCream(name: String, numCherries: Int, inCone: Boolean)

val sundae = LabelledGeneric[IceCream].
  to(IceCream("Sundae", 1, false))
// sundae: shapeless.::[String with shapeless.labelled.KeyTag[Symbol
  with shapeless.tag.Tagged[String("name")],String],shapeless.::[Int
  with shapeless.labelled.KeyTag[Symbol with shapeless.tag.Tagged[
  String("numCherries")],Int],shapeless.::[Boolean with shapeless.
  labelled.KeyTag[Symbol with shapeless.tag.Tagged[String("inCone")],
  Boolean],shapeless.HNil]] = Sundae :: 1 :: false :: HNil
```

Unlike the HList and Coproduct ops we have seen already, record ops syntax requires an explicit import from `shapeless.record`:

```
import shapeless.record._
```

6.4.1 Selecting fields

The `get` method from `shapeless.syntax.record` and its corresponding Selector type class allow us to fetch a field by tag.

```
sundae.get('name)
// res1: String = Sundae

sundae.get('numCherries)
// res2: Int = 1
```

Attempting to access an undefined field causes a compile error as we might expect:

```
sundae.get('nomCherries')
// <console>:22: error: No field Symbol with shapeless.tag.Tagged[String
  ("nomCherries")] in record shapeless.: [String with shapeless.
    labelled.KeyTag[Symbol with shapeless.tag.Tagged[String("name")],
    String], shapeless.: [Int with shapeless.labelled.KeyTag[Symbol with
    shapeless.tag.Tagged[String("numCherries")], Int], shapeless.: [
    Boolean with shapeless.labelled.KeyTag[Symbol with shapeless.tag.
    Tagged[String("inCone")], Boolean], shapeless.HNil]]
//      sundae.get('nomCherries)
//      ^
```

6.4.2 Updating and removing fields

The updated method and Updater type class allow us to modify fields by key. The remove method and Remover type class allow us to delete fields by key:

```
sundae.updated('numCherries', 3)
// res4: shapeless.: [String with shapeless.labelled.KeyTag[Symbol with
  shapeless.tag.Tagged[String("name")], String], shapeless.: [Int with
  shapeless.labelled.KeyTag[Symbol with shapeless.tag.Tagged[String("
  numCherries")], Int], shapeless.: [Boolean with shapeless.labelled.
  KeyTag[Symbol with shapeless.tag.Tagged[String("inCone")], Boolean],
  shapeless.HNil]] = Sundae :: 3 :: false :: HNil

sundae.remove('inCone')
// res5: (Boolean, shapeless.: [String with shapeless.labelled.KeyTag[
  Symbol with shapeless.tag.Tagged[String("name")], String], shapeless.
  .: [Int with shapeless.labelled.KeyTag[Symbol with shapeless.tag.
  Tagged[String("numCherries")], Int], shapeless.HNil]) = (false,
  Sundae :: 1 :: HNil)
```

The updateWith method and Modifier type class allow us to modify a field with an update function:

```
sundae.updateWith('name')("MASSIVE " + _)
// res6: shapeless.: [String with shapeless.labelled.KeyTag[Symbol with
  shapeless.tag.Tagged[String("name")], String], shapeless.: [Int with
  shapeless.labelled.KeyTag[Symbol with shapeless.tag.Tagged[String("
  numCherries")], Int], shapeless.: [Boolean with shapeless.labelled.
```



```
KeyTag[Symbol with shapeless.tag.Tagged[String("inCone")],Boolean],  
shapeless.HNil]] = MASSIVE Sundae :: 1 :: false :: HNil
```

6.4.3 Converting to a regular Map

The `toMap` method and `ToMap` type class allow us to convert a record to a Map:

```
sundae.toMap  
// res7: Map[Symbol with shapeless.tag.Tagged[_ >: String("inCone") with  
    String("numCherries") with String("name") <: String],Any] = Map('inCone -> false, 'numCherries -> 1, 'name -> Sundae)
```

6.4.4 Other operations

There are other record ops that we don't have room to cover here. We can rename fields, merge records, map over their values, and much more. See the source code of `shapeless.ops.record` and `shapeless.syntax.record` for more information.

6.5 Summary

In this chapter we explored a few of the type classes that are provided in the `shapeless.ops` package. We looked at `Last` and `Init` as two simple examples of the ops pattern, and built our own `Penultimate` and `Migration` type classes purely by chaining together existing building blocks.

Many of the ops type classes share a similar pattern. The easiest way to learn them is to look at the source code in `shapeless.ops` and `shapeless.syntax`.

In the next chapters we will look at two suites of that require further theoretical discussion. Chapter 7 discusses functional operations such as `map` and `flatMap` on `HLists`, and Chapter 8 discusses how to implement type classes that require type level representations of numbers. This knowledge will help us gain a complete understanding of nearly every type class that `shapeless.ops` has to offer.

Chapter 7

Functional operations on HLists

“Regular” Scala programs make heavy use of functional operations like `map` and `flatMap`. A question arises: can we perform similar operations on HLists? The answer is “yes”, but the heterogeneous element types require us to do things a little differently than in regular Scala. Unsurprisingly the mechanisms are type class based and there are a suite of ops type classes to help us out.

Before we delve in to the type classes themselves, we need to discuss how `shapeless` represents *polymorphic* functions suitable for mapping over heterogeneous data structures.

7.1 Motivation: mapping over an HList

We’ll motivate the discussion of polymorphic functions with a look at the `map` method. Figure 7.1 shows a type chart for mapping over a regular list. We start with a `List[A]`, supply a function `A => B`, and end up with a `List[B]`.

The heterogeneous element types in an `HList` cause this model to break down. Scala functions have fixed input and output types, so the result of our `map` will have to have the same element type in every position.

Ideally we’d like a `map` operation like the one shown in Figure 7.2, where the function inspects the type of each input and uses it to determine the type of




Figure 7.1: Mapping over a regular list (“monomorphic” map)

each output. This gives us a closed, composable transformation that retains the heterogeneous nature of the HList.




Figure 7.2: Mapping over a heterogeneous list (“polymorphic” map)

Unfortunately we can’t use Scala functions to implement this kind of operation. We need some new infrastructure.

7.2 Polymorphic functions

Shapeless represents polymorphic functions using a type called `Poly`. Here is a simplified explanation of how it works. Note that this isn’t real shapeless code—we’re eliding much that makes real shapeless `Polys` more flexible and easier to use.

7.2.1 How Polys work

At its core, a `Poly` is an object with a generic `apply` method. In addition to its regular parameter of type `A`, `Poly` accepts an implicit parameter of type `Case[A]`:

```
// This is not real shapeless code.
// It is purely for illustration.

trait Case[P, A] {
  type Result
  def apply(a: A): Result
}

trait Poly {
  def apply[A](arg: A)(implicit cse: Case[this.type, A]): cse.Result =
    cse.apply(arg)
}
```

When we define an actual `Poly`, we provide instances of `Case` for each parameter type we care about. These implement the actual function body:

```
// This is not real shapeless code.
// It is purely for illustration.

object myPoly extends Poly {
  implicit def intCase =
    new Case[this.type, Int] {
      type Result = Double
      def apply(num: Int): Double = num / 2.0
    }

  implicit def stringCase =
    new Case[this.type, String] {
      type Result = Int
      def apply(str: String): Int = str.length
    }
}
```

When we call `myPoly.apply`, the compiler searches for the relevant implicit `Case` and fills it in as usual:

```
myPoly.apply(123)
// res7: Double = 61.5
```

There is some subtle scoping behaviour here that allows the compiler to locate instances of `Case` without any additional imports. `Case` has an extra type

parameter `P` referencing the singleton type of the `Poly`. The implicit scope for `Case[P, A]` includes the companion objects for `Case`, `P`, and `A`. The companion object for `myPoly.type` is `myPoly` itself, so `Cases` defined in the body of the `Poly` are always in scope no matter where the call site is.

7.2.2 Poly syntax

The code above isn't real shapeless code. Fortunately, shapeless makes `Polys` much simpler to define. Here's our `myPoly` function rewritten in proper syntax:

```
import shapeless._

object myPoly extends Poly1 {
  implicit val intCase: Case.Aux[Int, Double] =
    at(num => num / 2.0)

  implicit val stringCase: Case.Aux[String, Int] =
    at(str => str.length)
}
```

There are a few key differences with our earlier toy syntax:

1. We're extending a trait called `Poly1` instead of `Poly`. Shapeless has a `Poly` type and a set of subtypes, `Poly1` through `Poly22`, supporting different arities of polymorphic function.
2. The `Case.Aux` types doesn't seem to reference the singleton type of the `Poly`. `Case.Aux` is actually actually a type alias defined within the body of `Poly1`. The singleton type is there—we just don't see it.
3. We're using a helper method, `at`, to define cases. This lets us eliminate a lot of boilerplate such as writing out full definitions of `Result` and `apply`.

Syntactic differences aside, the shapeless version of `myPoly` is functionally identical to our toy version. We can call it with an `Int` or `String` parameter and get back a result of the corresponding return type:

```
myPoly.apply(123)
// res9: myPoly.intCase.Result = 61.5

myPoly.apply("hello")
// res10: myPoly.stringCase.Result = 5
```

Shapeless also supports Polys with more than one parameter. Here is a binary example:

```
object multiply extends Poly2 {
  implicit val intIntCase: Case.Aux[Int, Int, Int] =
    at((a, b) => a * b)

  implicit val intStrCase: Case.Aux[Int, String, String] =
    at((a, b) => b.toString * a)
}

multiply(3, 4)
// res11: multiply.intIntCase.Result = 12

multiply(3, "4")
// res12: multiply.intStrCase.Result = 444
```

Because Cases are just implicit values, we can define cases based on type classes and do all of the advanced implicit resolution covered in previous chapters. Here's a simple example that totals numbers in different contexts:

```
import scala.math.Numeric

object total extends Poly1 {
  implicit def baseCase[A](implicit num: Numeric[A]): Case.Aux[A, Double] =
    at(num.toDouble)

  implicit def optionCase[A](implicit num: Numeric[A]): Case.Aux[Option[A], Double] =
    at(opt => opt.map(num.toDouble).getOrElse(0.0))

  implicit def listCase[A](implicit num: Numeric[A]): Case.Aux[List[A], Double] =
```

```

    at(list => num.toDouble(list.sum))
  }

total(10)
// res14: Double = 10.0

total(Option(20.0))
// res15: Double = 20.0

total(List(1L, 2L, 3L))
// res16: Double = 6.0

```

Idiosyncrasies of type inference

Poly pushes Scala's type inference out of its comfort zone. We can easily confuse the compiler by asking it to do too much inference at once. For example, the following code compiles ok:

```

val a = myPoly.apply(123)
val b: Double = a

```

However, if we try to combine the two lines in to one we get a compilation error:

```

val a: Double = myPoly.apply(123)
// <console>:17: error: type mismatch;
//   found   : Int(123)
//   required: myPoly.ProductCase.Aux[shapeless.HNil,?]
//   (which expands to)  shapeless.poly.Case[myPoly.type,
//                        shapeless.HNil]{type Result = ?}
//       val a: Double = myPoly.apply(123)
//                                ^

```

If we give the compiler a hint by telling it what the parameter type is, the code compiles again:

```

val a: Double = myPoly.apply[Int](123)
// a: Double = 61.5

```


This behaviour is confusing and annoying. Unfortunately there are no concrete rules to follow to avoid problems. The only general guideline is to avoid over-constraining the compiler—solve one constraint at a time and give it a hint when it gets stuck.

7.3 Mapping and flatMapping using Poly

Shapeless provides a suite of functional operations based on Poly, each implemented as an ops type class. Let's look at map and flatMap as examples. Here's map:

```
import shapeless._

object sizeOf extends Poly1 {
  implicit val intCase: Case.Aux[Int, Int] =
    at(identity)

  implicit val stringCase: Case.Aux[String, Int] =
    at(_.length)

  implicit val booleanCase: Case.Aux[Boolean, Int] =
    at(bool => if(bool) 1 else 0)
}

(10 :: "hello" :: true :: HNil).map(sizeOf)
// res1: shapeless.:::[Int,shapeless.:::[Int,shapeless.:::[Int,shapeless.
//      HNil]]] = 10 :: 5 :: 1 :: HNil
```

Note that the elements in the resulting HList have types matching the Cases in sizeOf. We can use map with any Poly that provides Cases for every member of our starting HList. If the compiler can't find a Case for a particular member, we get a compile error:

```
(1.5 :: HNil).map(sizeOf)
// <console>:17: error: could not find implicit value for parameter
//      mapper: shapeless.ops.hlist.Mapper[sizeOf.type,shapeless.:::[Double,
//      shapeless.HNil]]
//      (1.5 :: HNil).map(sizeOf)
```


7.4 Folding using Poly

In addition to `map` and `flatMap`, `shapeless` also provides `foldLeft` and `foldRight` operations based on `Poly2`:

```
import shapeless._

object sum extends Poly2 {
  implicit val intIntCase: Case.Aux[Int, Int, Int] =
    at((a, b) => a + b)

  implicit val intStringCase: Case.Aux[Int, String, Int] =
    at((a, b) => a + b.length)
}

(10 :: "hello" :: 100 :: HNil).foldLeft(0)(sum)
// res7: Int = 115
```

We can also `reduceLeft`, `reduceRight`, `foldMap`, and so on. Each operation has its own associated type class. We'll leave it as an exercise to the reader to investigate the available operations.

7.5 Defining type classes using Poly

We can use `Poly` and type classes like `Mapper` and `FlatMapper` as building blocks for our own type classes. As an example let's build a type class for mapping from one case class to another:

```
trait ProductMapper[A, B, P] {
  def apply(a: A): B
}
```

We can create an instance of `ProductMapper` using `Mapper` and a pair of `Generics`:

```

import shapeless._
import shapeless.ops.hlist

implicit def genericProductMapper[
  A, B,
  P <: Poly,
  ARepr <: HList,
  BRepr <: HList
](
  implicit
  aGen: Generic.Aux[A, ARepr],
  bGen: Generic.Aux[B, BRepr],
  mapper: hlist.Mapper.Aux[P, ARepr, BRepr]
): ProductMapper[A, B, P] =
  new ProductMapper[A, B, P] {
    def apply(a: A): B =
      bGen.from(mapper.apply(aGen.to(a)))
  }

```

Interestingly, although we define a type `P` for our `Poly`, we don't reference any values of type `P` anywhere in our code. The `Mapper` type class uses implicit resolution to find `Cases`, so we only need to know the type to do the mapping.

Let's create an extension method to make `ProductMapper` easier to use. We only want the user to specify the type of `B` at the call site, so we use some indirection to allow the compiler to infer the type of the `Poly` from a value parameter:

```

implicit class ProductMapperOps[A](a: A) {
  class Builder[B] {
    def apply[P <: Poly](poly: P)
      (implicit prodMap: ProductMapper[A, B, P]): B =
      prodMap(a)
  }

  def mapTo[B]: Builder[B] =
    new Builder[B]
}

```

Here's an example of the method's use:

```
object conversions extends Poly1 {
  implicit val intCase: Case.Aux[Int, Boolean] = at(_ > 0)
  implicit val boolCase: Case.Aux[Boolean, Int] = at(if(_) 1 else 0)
  implicit val strCase: Case.Aux[String, String] = at(identity)
}

case class IceCream1(name: String, numCherries: Int, inCone: Boolean)
case class IceCream2(name: String, hasCherries: Boolean, numCones: Int)

IceCream1("Sundae", 1, false).mapTo[IceCream2](conversions)
// res2: IceCream2 = IceCream2(Sundae,true,0)
```

The `mapTo` syntax looks like a single method call, but is actually two calls: one call to `mapTo` to fix the `B` type parameter, and one call to `Builder.apply` to specify the `Poly`. Some of shapeless' built-in ops extension methods use similar tricks.

7.6 Summary

In this chapter we discussed *polymorphic functions* whose return types vary based on the types of their parameters. We saw how shapeless' `Poly` type is defined, and how it is used to implement functional operations such as `map`, `flatMap`, `foldLeft`, and `foldRight`.

Each operation is implemented as an extension method on `HList`, based on a corresponding type class: `Mapper`, `FlatMapper`, `LeftFolder`, and so on. We can use these type classes, `Poly`, and the techniques from Section 4.3 to create our own type classes involving sequences of sophisticated transformations.

Chapter 8

Counting with types

From time to time we need to count things at the type level. For example, we may need to know the length of an `HList` or the number of terms we have expanded so far in a computation. We can represent numbers as values easily enough, but if we want to influence implicit resolution we need to represent them as the type level. This chapter covers the theory behind counting with types, and provides some compelling use cases for type class derivation.

8.1 Representing numbers as types

Shapeless uses “Church encoding” to represent natural numbers at the type level. It provides a type `Nat` with two subtypes: `_0` representing zero, and `Succ[N]` representing `N+1`:

```
import shapeless.{Nat, Succ}

type Zero = Nat._0
type One  = Succ[Zero]
type Two  = Succ[One]
// etc...
```

Shapeless provides aliases for the first 22 `Nat`s as `Nat._N`:

```
Nat._1
Nat._2
Nat._3
// etc...
```

`Nat` has no runtime semantics. We have to use the `ToInt` type class to convert a `Nat` to a runtime `Int`:

```
import shapeless.ops.nat.ToInt

val toInt = ToInt[Two]

toInt.apply()
// res7: Int = 2
```

The `Nat.toInt` method provides a convenient shorthand for calling `nat.apply()`. It accepts the instance of `ToInt` as an implicit parameter:

```
Nat.toInt[Nat._3]
// res8: Int = 3
```

8.2 Length of generic representations

One use case for `Nat` is determining the length of `HLists` and `Coproducts`. `Shapeless` provides the `shapeless.ops.hlist.Length` and `shapeless.ops.coproduct.Length` type classes for this.

Because of the similarity of the names, we typically import the `hlist` and `coproduct` packages refer to the relevant type classes as `package.Length`:

```
import shapeless._
import shapeless.ops.nat.ToInt
import shapeless.ops.{hlist, coproduct}

val hlistLength = hlist.Length[String :: Int :: Boolean :: HNil]
// hlistLength: shapeless.ops.hlist.Length[shapeless.::[String,shapeless.::[Int,shapeless.::[Boolean,shapeless.HNil]]]]{type Out =
  shapeless.Succ[shapeless.Succ[shapeless.Succ[shapeless._0]]]} =
```



```

    shapeless.ops.hlist$Length$$anon$3@73551364

val coproductLength = coproduct.Length[Double :+: Char :+: CNil]
// coproductLength: shapeless.ops.coproduct.Length[shapeless.+:[Double,
    shapeless.+:[Char,shapeless.CNil]]]{type Out = shapeless.Succ[
    shapeless.Succ[shapeless._0]]} = shapeless.ops.
    coproduct$Length$$anon$29@2455fe3d

```

Instances of `Length` have a type member `Out` that represents the length as a `Nat`. We can either summon an instance of `ToInt` ourselves:

```

ToInt[hlistLength.Out].apply()
// res0: Int = 3

```

or use the `Nat.toInt` helper:

```

Nat.toInt[coproductLength.Out]
// res1: Int = 2

```

Let's use this in a concrete example. We'll create a `SizeOf` type class that counts the number of fields in a case class and exposes it as a simple `Int`:

```

trait SizeOf[A] {
  def value: Int
}

def sizeOf[A](implicit size: SizeOf[A]): Int =
  size.value

```

To create an instance of `SizeOf` we need three things:

1. a `Generic` to calculate the corresponding `HList` type;
2. a `Length` to calculate the length of the `HList` as a `Nat`;
3. a `ToInt` to convert the `Nat` to an `Int`.

Here's a working implementation written in the style described in Chapter 4:

```
implicit def genericSizeOf[A, L <: HList, N <: Nat](
  implicit
    generic: Generic.Aux[A, L],
    size: hlist.Length.Aux[L, N],
    sizeToInt: ToInt[N]
): SizeOf[A] =
  new SizeOf[A] {
    val value = sizeToInt.apply()
  }
```

We can test our code as follows:

```
case class IceCream(name: String, numCherries: Int, inCone: Boolean)

sizeOf[IceCream]
// res3: Int = 3
```

8.3 Case study: random value generator

Property-based testing libraries like `ScalaCheck`¹ use type classes to generate random data for unit tests. For example, `ScalaCheck` provides the `Arbitrary` type class that we can use as follows:

```
import org.scalacheck._

for(i <- 1 to 3) println(Arbitrary.arbitrary[Int].sample)
// Some(-2083372342)
// Some(-2147483648)
// Some(1579167812)

for(i <- 1 to 3) println(Arbitrary.arbitrary[(Boolean, Byte)].sample)
// Some((true,0))
// Some((true,-29))
// Some((true,127))
```

`ScalaCheck` provides built-in instances of `Arbitrary` for a wide range of standard Scala types. However, creating instances of `Arbitrary` for user ADTs is

¹<https://scalacheck.org/>

still a time-consuming manual process. This makes shapeless integration via libraries like `scalacheck-shapeless`² very attractive.

In this section we will create a simple `Random` type class to generate random values of user-defined ADTs. We will show how `Length` and `Nat` form a crucial part of the implementation. As usual we start with the definition of the type class itself:

```
trait Random[A] {  
  def get: A  
}  
  
def random[A](implicit r: Random[A]): A =  
  r.get
```

8.3.1 Simple random values

Let's start with some basic instances of `Random`:

```
// Helper method for creating instances:  
def createRandom[A](func: () => A): Random[A] =  
  new Random[A] {  
    def get = func()  
  }  
  
// Random numbers from 0 to 9:  
implicit val intRandom: Random[Int] =  
  createRandom(() => scala.util.Random.nextInt(10))  
  
// Random characters from A to Z:  
implicit val charRandom: Random[Char] =  
  createRandom(() => ('A'.toInt + scala.util.Random.nextInt(26)).toChar)  
  
// Random booleans:  
implicit val booleanRandom: Random[Boolean] =  
  createRandom(() => scala.util.Random.nextBoolean)
```

We can use these simple generators via the `random` method as follows:

²<https://github.com/alexarchambault/scalacheck-shapeless>

```

for(i <- 1 to 3) println(random[Int])
// 0
// 8
// 9

for(i <- 1 to 3) println(random[Char])
// V
// N
// J

```

8.3.2 Random products

We can create random values for products using the `Generic` and `HList` techniques from Chapter 3:

```

import shapeless._

implicit def genericRandom[A, R](
  implicit
    gen: Generic.Aux[A, R],
    random: Lazy[Random[R]]
): Random[A] =
  createRandom(() => gen.from(random.value.get))

implicit val hnilRandom: Random[HNil] =
  createRandom(() => HNil)

implicit def hlistRandom[H, T <: HList](
  implicit
    hRandom: Random[H],
    tRandom: Lazy[Random[T]]
): Random[H :: T] =
  createRandom(() => hRandom.get :: tRandom.value.get)

```

This gets us as far as summoning random instances for case classes:

```

case class Cell(col: Char, row: Int)

for(i <- 1 to 5) println(random[Cell])
// Cell(H,1)

```

```
// Cell(D,4)
// Cell(D,7)
// Cell(V,2)
// Cell(R,4)
```

8.3.3 Random coproducts

This is where we start hitting problems. Generating a random instance of a coproduct involves choosing a random subtype. Let's start with a naïve implementation:

```
implicit val cnilRandom: Random[CNil] =
  createRandom(() => throw new Exception("Inconceivable!"))
// cnilRandom: Random[shapeless.CNil] = $anon$1@36616c2

implicit def coproductRandom[H, T <: Coproduct](
  implicit
    hRandom: Random[H],
    tRandom: Lazy[Random[T]]
): Random[H :+: T] =
  createRandom { () =>
    val chooseH = scala.util.Random.nextDouble < 0.5
    if(chooseH) Inl(hRandom.get) else Inr(tRandom.value.get)
  }
// coproductRandom: [H, T <: shapeless.Coproduct](implicit hRandom:
  Random[H], implicit tRandom: shapeless.Lazy[Random[T]])Random[
  shapeless :+: [H,T]]
```

There problems with this implementation lie in the 50/50 choice in calculating chooseH. This creates an uneven probability distribution. For example, consider the following type:

```
sealed trait Light
case object Red extends Light
case object Amber extends Light
case object Green extends Light
```

The Repr for Light is Red :+: Amber :+: Green :+: CNil. An instance of Random for this type will choose Red 50% of the time and Amber :+: Green

`:+:` `CNil` 50% of the time. A correct distribution would be 33% Red and 67% Amber `:+:` Green `:+:` `CNil`.

And that's not all. If we look at the overall probability distribution we see something even more alarming:

- Red is chosen 50% of the time
- Amber is chosen 25% of the time
- Green is chosen 12.5% of the time
- `CNil` is chosen 6.75% of the time

Our coproduct instances will throw exceptions 6.75% of the time!

```
for(i <- 1 to 100) random[Light]
// java.lang.Exception: Inconceivable!
// ...
```

To fix this problem we have to alter the probability of choosing H over T. The correct behaviour should be to choose H $1/n^{\text{th}}$ of the time, where n is the length of the coproduct. This ensures an even probability distribution across the subtypes of the coproduct, which means we never call `cnilProduct.get`. Here's an updated implementation:

```
import shapeless.ops.coproduct
import shapeless.ops.nat.ToInt

implicit def coproductRandom[H, T <: Coproduct, L <: Nat]({
  implicit
    hRandom: Random[H],
    tRandom: Lazy[Random[T]],
    tLength: coproduct.Length.Aux[T, L],
    tLengthAsInt: ToInt[L]
}): Random[H :+: T] = {
  createRandom { () =>
    val length = 1 + tLengthAsInt()
    val chooseH = scala.util.Random.nextDouble < (1.0 / length)
    if(chooseH) Inl(hRandom.get) else Inr(tRandom.value.get)
```

```
}
}
```

With these modifications we can generate random values of any product or coproduct:

```
for(i <- 1 to 5) println(random[Light])
// Green
// Red
// Red
// Red
// Green
```

Generating test data for ScalaCheck tests normally requires a great deal of boilerplate. This is a particularly compelling use case for shapeless.

8.4 Other operations involving Nat

Shapeless provides a suite of other operations based on Nat. The apply methods on HList and Coproduct and can accept Nats as value or type parameters:

```
import shapeless._

val hlist = 123 :: "foo" :: true :: 'x' :: HNil

hlist.apply[Nat._1]
// res1: String = foo

hlist.apply[Nat._3]
// res2: Char = x
```

There are also operations such as take, drop, slice, and updatedAt:

```
hlist.take(Nat._3).drop(Nat._1)
// res3: shapeless.::[String,shapeless.::[Boolean,shapeless.HNil]] = foo
//      :: true :: HNil
```

```
hlist.updatedAt(Nat._1, "bar").updatedAt(Nat._2, "baz")
// res4: shapeless.::[Int,shapeless.::[String,shapeless.::[String,
    shapeless.::[Char,shapeless.HNil]]]] = 123 :: bar :: baz :: x ::
    HNil
```

These operations and their associated type classes are useful for manipulating individual elements within a product or coproduct.

8.5 Summary

In this chapter we discussed how `shapeless` represents natural numbers and how we can use them in type classes. We saw some predefined ops type classes that let us do things like calculate lengths and access elements by index, and created our own type classes that use `Nat` in other ways.

Between `Nat`, `Poly`, and the variety of type classes and examples we have seen in Part II, we have seen just a small fraction of the toolbox provided in `shapeless.ops`. There are many other ops type classes that provide a comprehensive foundation on which to build our own code.

Preparing for launch

With Part II's look at `shapeless.ops` we have arrived at the end of this guide. We hope you found it useful for understanding this fascinating and powerful library, and wish you all the best on your future journeys as a type astronaut.

As functional programmers we value abstraction above all else. Concepts like functors and monads arise from years of programming research: writing code, spotting patterns, and making abstractions to remove redundancy. Shapeless raises the bar for abstraction in Scala. Tools like `Generic` and `LabelledGeneric` provide an interface for abstracting over data types that were previously frustratingly unique and distinct.

There have traditionally been two barriers to entry for aspiring new shapeless users. The first is the wealth of theoretical knowledge and implementation detail required to understand the patterns we need. Hopefully this guide has helped in this regard.

The second barrier is the fear and uncertainty surrounding a library that is seen as “academic” or “advanced”. We can overcome this by sharing knowledge—use cases, pros and cons, implementation strategies, and so on—to widen the understanding of this valuable tool. So please share this book with a friend... and let's scrap some boilerplate together!

TODOs

Still to do:

- Mention type class instance selection in Chapter 4
 - Low and high priority
 - `LowPriorityImplicit`
- Function interop
- Performance
 - `cachedImplicit`
 - Maybe Cached
 - Maybe
- Check cross references
- **DONE** Generic applied to tuples
- **DONE** Mention SI-7046 in Chapter 3
- **DONE** Complete the “debugging” section in Chapter 4
- **DONE** Built-in `HList` and `Coproduct` operations
 - **DONE** Migrating case class as a basic example
- **DONE** Polymorphic functions
 - **DONE** Mapping over `HLists` as an example
- **DONE** Counting with `Nat`
 - **DONE** Generating Arbitrary instances as an example

- **DONE** Callout box on quirkiness of type inference with poly:
 - **DONE** `val len1: Int = lengthPoly("foo")` fails, but...
 - **DONE** `val len2 = lengthPoly("foo")` compiles, but...
 - **DONE** `val len3: Int = lengthPoly[String]("foo")` fails
- **DONE** Built-in record operations
- **DONE** Final summary
- **SHIP IT!**